

RISE in MLIR

A functional Pattern-based Dialect

Martin Lücke | Michel Steuwer | Aaron Smith



THE UNIVERSITY
of EDINBURGH



University
of Glasgow

supported by a Google Faculty Award sponsored by Jacques Pienaar and Albert Cohen

Why RISE?

Machine Learning Systems are Stuck in a Rut

[HotOS'19]

Paul Barham
Google Brain

Michael Isard
Google Brain

Abstract

In this paper we argue that systems for numerical computing are stuck in a local basin of performance and programmability. Systems researchers are doing an excellent job improving the performance of 5-year-old benchmarks, but gradually making it harder to explore innovative machine learning research ideas.

We explain how the evolution of hardware accelerators favors compiler back ends that hyper-optimize large monolithic kernels, show how this reliance on high-performance but inflexible kernels reinforces the dominant style of programming model, and argue these programming abstractions lack expressiveness, maintainability, and modularity; all of which hinders research progress.

We conclude by noting promising directions in the field, and advocate steps to advance progress towards high-performance general purpose numerical computing systems on modern accelerators.

ACM Reference Format:

Paul Barham and Michael Isard. 2019. Machine Learning Systems are Stuck in a Rut. In *Workshop on Hot Topics in Operating Systems (HotOS '19)*, May 13–15, 2019, Bertinoro, Italy. ACM, New York, NY, USA, 7 pages. <https://doi.org/10.1145/3317550.3321441>




Figure 1. Conv2D operation with 3x3 kernel, stride=2

with 16 times fewer training parameters than the convolutional neural network (CNN) we were comparing it to, implementations in both TensorFlow[2] and PyTorch[3] were much slower and ran out of memory with much smaller models. We wanted to understand why.

1.1 New ideas often require new primitives

We won't discuss the full details of Capsule networks in this paper¹, but for our purposes it is sufficient to consider a simplified form of the inner loop, which is

Original authors
of TensorFlow



Why RISE?

Machine Learning Systems are Stuck in a Rut

Paul Barham
Google Brain

Michael Isard
Google Brain

Abstract

In this paper we argue that systems for numerical computing are stuck in a local basin of performance and programmability. Systems researchers are doing an excellent job improving the performance of 5-year-old benchmarks, but gradually making it harder to explore innovative machine learning research ideas.

We explain how the evolution of hardware accelerators favors compiler back ends that **hyper-optimize large monolithic kernels**, show how this reliance on high-performance but **inflexible** kernels reinforces the dominant style of programming model, and argue these programming abstractions lack expressiveness, maintainability, and **modularity**; all of which **hinders research progress**.

We conclude by noting promising directions in the field, and advocate steps to advance progress towards high-performance general purpose numerical computing systems on modern accelerators.

ACM Reference Format:

Paul Barham and Michael Isard. 2019. Machine Learning Systems are Stuck in a Rut. In *Workshop on Hot Topics in Operating Systems (HotOS '19)*, May 13–15, 2019, Bertinoro, Italy. ACM, New York, NY, USA, 7 pages. <https://doi.org/10.1145/3317550.3321441>




Figure 1. Conv2D operation with 3×3 kernel, stride=2

with 16 times fewer training parameters than the convolutional neural network (CNN) we were comparing it to, implementations in both TensorFlow[2] and PyTorch[3] were much slower and ran out of memory with much smaller models. We wanted to understand why.

1.1 New ideas often require new primitives

We won't discuss the full details of Capsule networks in this paper¹, but for our purposes it is sufficient to consider a simplified form of the inner loop, which is

Why RISE?

Machine Learning Systems are Stuck in a Rut

Paul Barham
Google Brain

Michael Isard
Google Brain

Abstract

In this paper we argue that systems for numerical computing are stuck in a local basin of performance and programmability. Systems researchers are doing an excellent job improving the performance of 5-year-old benchmarks, but gradually making it harder to explore innovative machine learning research ideas.

We explain how the evolution of hardware accelerators favors compiler back ends that **hyper-optimize large monolithic kernels**, show how this reliance on high-performance but **inflexible** kernels reinforces the dominant style of programming model, and argue these programming abstractions lack expressiveness, maintainability, and **modularity**; all of which **hinders research progress**.

We conclude by noting promising directions in the field and advocate steps to advance progress towards




Figure 1. Conv2D operation with 3x3 kernel, stride=2 with 16 times fewer training parameters than the convo-

We should aim for more principled higher level intermediate representations

Paul Barham and Michael Isard. 2019. Machine Learning Systems are Stuck in a Rut. In *Workshop on Hot Topics in Operating Systems (HotOS '19)*, May 13–15, 2019, Bertinoro, Italy. ACM, New York, NY, USA, 7 pages. <https://doi.org/10.1145/3317550.3321441>

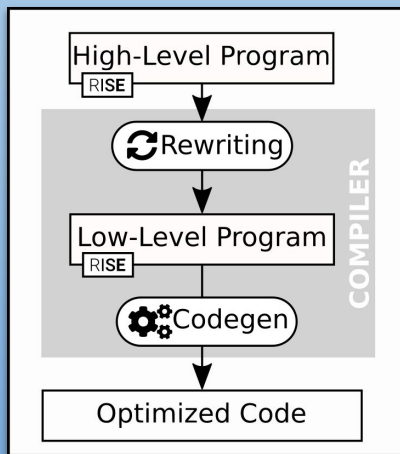
smaller models. we wanted to understand why.

1.1 New ideas often require new primitives

We won't discuss the full details of Capsule networks in this paper¹, but for our purposes it is sufficient to consider a simplified form of the inner loop, which is


RISE - A functional pattern-based data-parallel language

- **RISE** (<https://rise-lang.org/>) is a spiritual successor to the Lift project
- Computations represented using compositions of flexible and generic patterns
- Rewriting system to explore optimization choices




RISE by example: Matrix Multiplication

```
fun(A : N.K.float ⇒ fun(B : K.M.float ⇒  
  A ▷ map(fun(arow ⇒  
    B ▷ transpose ▷ map(fun(bcol ⇒  
      zip(arow, bcol) ▷ map(*) ▷ reduce(+, 0) )) )) ))
```



RISE by example: Matrix Multiplication

```
fun(A : N.K.float ⇒ fun(B : K.M.float ⇒  
  A ▷ map(fun(arow ⇒  
    B ▷ transpose ▷ map(fun(bcol ⇒  
      zip(arow, bcol) ▷ map(*) ▷ reduce(+, 0)) )) )) ))
```





dot product computation:

$$\sum arow_i * bcol_i$$

RISE by example: Matrix Multiplication

```
fun(A : N.K.float ⇒ fun(B : K.M.float ⇒  
  A ▷ map(fun(arow ⇒  
    B ▷ transpose ▷ map(fun(bcol ⇒  
      zip(arow, bcol) ▷ map(*) ▷ reduce(+, 0) )) )) ))
```



Optimization choices for Matrix Multiplication



```
fun(A : N.K.float ⇒ fun(B : K.M.float ⇒
  A ▷ map(fun(arow ⇒
    B ▷ transpose ▷ map(fun(bcol ⇒
      zip(arow, bcol) ▷ map(*) ▷ reduce(+, 0) )) )) ))
```



RISE compiler

[illegible][illegible]

```
kernel void mm(global float4* const A,
               global float4* const B,
               global float4*2 C, uint n) {
    uint i = get_global_id(0);
    uint j = get_global_id(1);
    uint nv4 = n >> 2;
    float4 ab = (float4)(0.f);
    for (uint k = 0; k < nv4; ++k) {
        float4 a0 = A[2*i + k];
        float4 a1 = A[2*i+1 + k];
        float4 b0 = B[2*j + k];
        float4 b1 = B[2*j+1 + k];
        ab += (float4)(dot(a0, b0), dot(a0, b1),
                       dot(a1, b0), dot(a1, b1));
    }
    C[i] = 2*i*(nv4+1) + j;
    C[j] = ab;
    C[i+j + (nv4+1)] = ab.dot2;
}
```



Optimization choices for Matrix Multiplication

```
fun(A : N.K.float ⇒ fun(B : K.M.float ⇒
  A ▷ map(fun(arow ⇒
    B ▷ transpose ▷ map(fun(bcol ⇒
      zip(arow, bcol) ▷ map(*) ▷ reduce(+, 0) )) )) ))
```





RISE compiler



```
kernel void mm(global float4* const A,
               global float4* const B,
               global float4* const C,
               uint n, uint m, uint k) {
    for (uint i = 0; i < n; i++) {
        for (uint j = 0; j < m; j++) {
            float4 sum = 0;
            for (uint k = 0; k < m; k++) {
                sum += A[i * m + k] * B[k * m + j];
            }
            C[i * m + j] = sum;
        }
    }
}
```

Rewrite Rules

$\text{map}(f, A) \mapsto \text{join}(\text{map}(\text{map}(f), \text{split}(n, A)))$



RISE Compilation

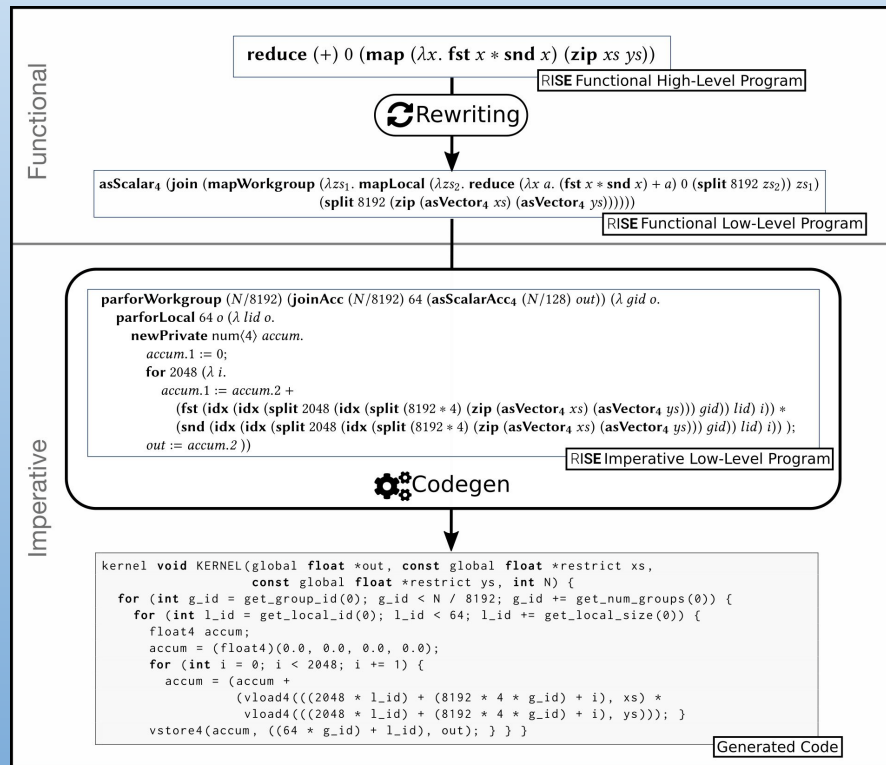
Compilation is divided into two steps:

1. Translation from functional to imperative low level program
2. Produce generated code

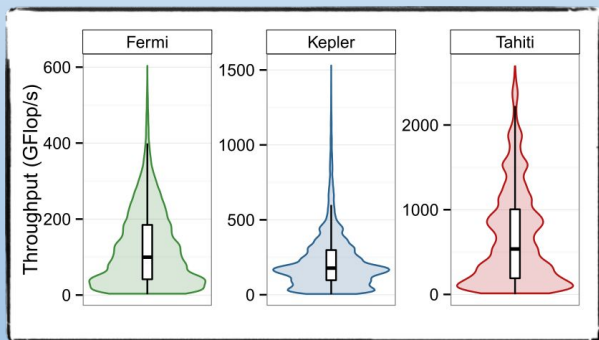
Compilation described formally in:

Strategy Preserving Compilation for Parallel Functional Code by Robert Atkey, Michel Steuwer, Sam Lindley, and Christophe Dubach.

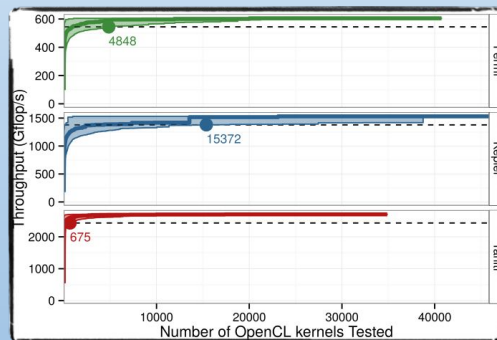
<https://arxiv.org/abs/1710.08332>




High Performance Results for Matrix Multiplication



Only few generated code with very good performance



Still: One can expect to find a good performing kernel quickly!



Performance close or better than hand-tuned library code

[GPGPU'16]

RISE - The Caveats

- Does everyone have to write functional programs now?
- Academic work written in Scala
- Does not integrate well with existing compiler infrastructures



MLIR - Multi-Level Intermediate Representation

- Extensible infrastructure to define compiler intermediate representations
- Dialects can capture different levels of abstraction:
 - High-level domain specific ----- Hardware specific backend
- Existing dialects available for:
 - TensorFlow / TensorFlow Lite
 - Targeting GPUs
 - ...
 - Performing polyhedral optimizations
 - LLVM IR



MLIR - Martin Lücke Intermediate Representation

- Extensible infrastructure to define compiler intermediate representations
- Dialects can capture different levels of abstraction:
 - High-level domain specific ----- Hardware specific backend
- Existing dialects available for:
 - TensorFlow / TensorFlow Lite
 - Targeting GPUs
 - ...
 - Performing polyhedral optimizations
 - LLVM IR



MLIR - Multi-Level Intermediate Representation

- Extensible infrastructure to define compiler intermediate representations
- Dialects can capture different levels of abstraction:
 - High-level domain specific ----- Hardware specific backend
- Existing dialects available for:
 - TensorFlow / TensorFlow Lite
 - Targeting GPUs
 - ...
 - Performing polyhedral optimizations
 - LLVM IR

RISE in MLIR = Lambda Calculus + Patterns in MLIR

- **RISE** in MLIR opens up opportunities to integrate with other MLIR dialects
- We do not have to write programs in **RISE** directly, but lower from domain-specific dialects to it
- MLIR is written in C++, using an established toolchain -> widely usable
- Natural integration with existing toolchains `#include libMLIR`

-> to implement **RISE** as an MLIR dialect we implement:

- λ -calculus
- patterns

RISE dialect by example: Matrix Multiplication

```
1 func @mm(%out:memref<1024x1024xf32>, %inA:memref<1024x1024xf32>, %inB:memref<1024x1024xf32>) {
2   %A = rise.in %inA : !rise.array<1024, array<1024, scalar<f32>>>
3   %B = rise.in %inB : !rise.array<1024, array<1024, scalar<f32>>>
4   %f1 = rise.lambda (%arow : !rise.array<1024, scalar<f32>>) → !rise.array<1024, scalar<f32>> {
5     %f2 = rise.lambda (%bcol : !rise.array<1024, scalar<f32>>) → !rise.array<1024, scalar<f32>> {
6       %zip = rise.zip #rise.nat<1024> #rise.scalar<f32> #rise.scalar<f32>
7       %zipped = rise.apply %zip, %arow, %bcol
8       %f = rise.lambda (%tuple : !rise.tuple<scalar<f32>, scalar<f32>>) → !rise.scalar<f32> {
9         %fstFun = rise.fst #rise.scalar<f32> #rise.scalar<f32>
10        %sndFun = rise.snd #rise.scalar<f32> #rise.scalar<f32>
11        %fst = rise.apply %fstFun, %tuple
12        %snd = rise.apply %sndFun, %tuple
13        %result = rise.embed(%fst, %snd) {
14          %res = mulf %fst, %snd : f32
15          return %res : f32
16        } : !rise.scalar<f32>
17        rise.return %result : !rise.scalar<f32>
18      }
19      %map = rise.mapSeq #rise.nat<1024> #rise.tuple<scalar<f32>, scalar<f32>> #rise.scalar<f32>
20      %multipliedArray = rise.apply %map, %f, %zipped
21      %add = rise.lambda (%a : !rise.scalar<f32>, %b : !rise.scalar<f32>) → !rise.scalar<f32>> {
22        %result = rise.embed(%a, %b) {
23          %res = addf %a, %b : f32
24          return %res : f32
25        } : !rise.scalar<f32>
26        rise.return %result : !rise.scalar<f32>
27      }
28      %init = rise.literal #rise.lit<0.0>
29      %reduce = rise.reduceSeq #rise.nat<1024> #rise.scalar<f32> #rise.scalar<f32>
30      %result = rise.apply %reduce, %add, %init, %multipliedArray
31      rise.return %result : !rise.scalar<f32>
32    }
33    %mapB = rise.mapSeq #rise.nat<1024> #rise.array<1024, scalar<f32>> #rise.array<1024, scalar<f32>>
34    %result = rise.apply %mapB, %f2, %B
35    rise.return %result : !rise.array<1024, array<1024, scalar<f32>>>
36  }
37  %mapA = rise.mapSeq #rise.nat<1024> #rise.array<1024, scalar<f32>> #rise.array<1024, scalar<f32>>
38  %result = rise.apply %mapA, %f1, %A
39  rise.out %out ← %result
40  return
41 }
```

RISE dialect by example: Matrix Multiplication

```
1 func @mm(%out:memref<1024x1024xf32>, %inA:memref<1024x1024xf32>, %inB:memref<1024x1024xf32>) {
2   %A = rise.in %inA : !rise.array<1024, array<1024, scalar<f32>>>
3   %B = rise.in %inB : !rise.array<1024, array<1024, scalar<f32>>>
4   %f1 = rise.lambda (%arow : !rise.array<1024, scalar<f32>>) → !rise.array<1024, scalar<f32>> {
5     %f2 = rise.lambda (%bcol : !rise.array<1024, scalar<f32>>) → !rise.array<1024, scalar<f32>> {
6       %zip = rise.zip #rise.nat<1024> #rise.scalar<f32> #rise.scalar<f32>
7       %zipped = rise.apply %zip, %arow, %bcol
8       %f = rise.lambda (%tuple : !rise.tuple<scalar<f32>, scalar<f32>>) → !rise.scalar<f32> {
9         %fstFun = rise.fst #rise.scalar<f32> #rise.scalar<f32>
10        %sndFun = rise.snd #rise.scalar<f32> #rise.scalar<f32>
11        %fst = rise.apply %fstFun, %tuple
12        %snd = rise.apply %sndFun, %tuple
13        %result = rise.embed(%fst, %snd) {
14          %res = mulf %fst, %snd : f32
15          return %res : f32
16        } : !rise.scalar<f32>
17        rise.return %result : !rise.scalar<f32>
18      }
19      %map = rise.mapSeq #rise.nat<1024> #rise.tuple<scalar<f32>, scalar<f32>> #rise.scalar<f32>
20      %multipliedArray = rise.apply %map, %f, %zipped
21      %add = rise.lambda (%a : !rise.scalar<f32>, %b : !rise.scalar<f32>) → !rise.scalar<f32>> {
22        %result = rise.embed(%a, %b) {
23          %res = addf %a, %b : f32
24          return %res : f32
25        } : !rise.scalar<f32>
26        rise.return %result : !rise.scalar<f32>
27      }
28      %init = rise.literal #rise.lit<0.0>
29      %reduce = rise.reduceSeq #rise.nat<1024> #rise.scalar<f32> #rise.scalar<f32>
30      %result = rise.apply %reduce, %add, %init, %multipliedArray
31      rise.return %result : !rise.scalar<f32>
32    }
33    %mapB = rise.mapSeq #rise.nat<1024> #rise.array<1024, scalar<f32>> #rise.array<1024, scalar<f32>>
34    %result = rise.apply %mapB, %f2, %B
35    rise.return %result : !rise.array<1024, array<1024, scalar<f32>>>
36  }
37  %mapA = rise.mapSeq #rise.nat<1024> #rise.array<1024, scalar<f32>> #rise.array<1024, scalar<f32>>
38  %result = rise.apply %mapA, %f1, %A
39  rise.out %out ← %result
40  return
41 }
```

Data Types

Function Types

Types

Data Types

```
2  %A = rise.in %inA : !rise.array<1024, array<1024, scalar<f32>>>
```

- Rise data types: array types, tuple types, scalar types
- Nested array types represent higher dimensional data
- Scalar wraps an arbitrary scalar MLIR type

Function Types

```
8  %f = rise.lambda (%t : !rise.tuple<scalar<f32>, scalar<f32>>) → scalar<f32>  
   // %f : !rise.fun<tuple<scalar<f32>, scalar<f32>> → scalar<f32>>
```

- Rise function types: types of lambda expressions and functional patterns
- Type system prevents mixing of function and data types

Patterns

```
1 func @mm(%out:memref<1024×1024xf32>, %inA:memref<1024×1024xf32>, %inB:memref<1024×1024xf32>) {
2   %A = rise.in %inA : !rise.array<1024, array<1024, scalar<f32>>>
3   %B = rise.in %inB : !rise.array<1024, array<1024, scalar<f32>>>
4   %f1 = rise.lambda (%arow : !rise.array<1024, scalar<f32>>) → !rise.array<1024, scalar<f32>> {
5     %f2 = rise.lambda (%bcol : !rise.array<1024, scalar<f32>>) → !rise.array<1024, scalar<f32>> {
6       %zip = rise.zip #rise.nat<1024> #rise.scalar<f32> #rise.scalar<f32>
7       %zipped = rise.apply %zip, %arow, %bcol
8       %f = rise.lambda (%tuple : !rise.tuple<scalar<f32>, scalar<f32>>) → !rise.scalar<f32> {
9         %fstFun = rise.fst #rise.scalar<f32> #rise.scalar<f32>
10        %sndFun = rise.snd #rise.scalar<f32> #rise.scalar<f32>
11        %fst = rise.apply %fstFun, %tuple
12        %snd = rise.apply %sndFun, %tuple
13        %result = rise.embed(%fst, %snd) {
14          %res = mulf %fst, %snd : f32
15          return %res : f32
16        } : !rise.scalar<f32>
17        rise.return %result : !rise.scalar<f32>
18      }
19      %map = rise.mapSeq #rise.nat<1024> #rise.tuple<scalar<f32>, scalar<f32>> #rise.scalar<f32>
20      %multipliedArray = rise.apply %map, %f, %zipped
21      %add = rise.lambda (%a : !rise.scalar<f32>, %b : !rise.scalar<f32>) → !rise.scalar<f32> {
22        %result = rise.embed(%a, %b) {
23          %res = addf %a, %b : f32
24          return %res : f32
25        } : !rise.scalar<f32>
26        rise.return %result : !rise.scalar<f32>
27      }
28      %init = rise.literal #rise.lit<0.0>
29      %reduce = rise.reduceSeq #rise.nat<1024> #rise.scalar<f32> #rise.scalar<f32>
30      %result = rise.apply %reduce, %add, %init, %multipliedArray
31      rise.return %result : !rise.scalar<f32>
32    }
33  }
34  %mapB = rise.mapSeq #rise.nat<1024> #rise.array<1024, scalar<f32>> #rise.array<1024, scalar<f32>>
35  %result = rise.apply %mapB, %f2, %B
36  rise.return %result : !rise.array<1024, array<1024, scalar<f32>>>
37 }
38 %mapA = rise.mapSeq #rise.nat<1024> #rise.array<1024, scalar<f32>> #rise.array<1024, scalar<f32>>
39 %result = rise.apply %mapA, %f1, %A
40 rise.out %out ← %result
41 }
```

Patterns

```
1 func @mm(%out:memref<1024×1024xf32>, %inA:memref<1024×1024xf32>, %inB:memref<1024×1024xf32>) {
2   %A = rise.in %inA : !rise.array<1024, array<1024, scalar<f32>>>
3   %B = rise.in %inB : !rise.array<1024, array<1024, scalar<f32>>>
4   %f1 = rise.lambda (%arow : !rise.array<1024, scalar<f32>>) → !rise.array<1024, scalar<f32>> {
5     %f2 = rise.lambda (%bcol : !rise.array<1024, scalar<f32>>) → !rise.array<1024, scalar<f32>> {
6       %zip = rise.zip #rise.nat<1024> #rise.scalar<f32> #rise.scalar<f32>
7       %zipped = rise.apply %zip, %arow, %bcol
8       %f = rise.lambda (%tuple : !rise.tuple<scalar<f32>, scalar<f32>>) → !rise.scalar<f32> {
9         %fstFun = rise.fst #rise.scalar<f32> #rise.scalar<f32>
10        %sndFun = rise.snd #rise.scalar<f32> #rise.scalar<f32>
11        %fst = rise.apply %fstFun, %tuple
12        %snd = rise.apply %sndFun, %tuple
13        %result = rise.embed(%fst, %snd) {
14          %res = mulf %fst, %snd : f32
15          return %res : f32
16        } : !rise.scalar<f32>
17        rise.return %result : !rise.scalar<f32>
18      }
19      %map = rise.mapSeq #rise.nat<1024> #rise.tuple<scalar<f32>, scalar<f32>> #rise.scalar<f32>
20      %multipliedArray = rise.apply %map, %f, %zipped
21      %add = rise.lambda (%a : !rise.scalar<f32>, %b : !rise.scalar<f32>) → !rise.scalar<f32> {
22        %result = rise.embed(%a, %b) {
23          %res = addf %a, %b : f32
24          return %res : f32
25        } : !rise.scalar<f32>
26        rise.return %result : !rise.scalar<f32>
27      }
28      %init = rise.literal #rise.lit<0.0>
29      %reduce = rise.reduceSeq #rise.nat<1024> #rise.scalar<f32> #rise.scalar<f32>
30      %result = rise.apply %reduce, %add, %init, %multipliedArray
31      rise.return %result : !rise.scalar<f32>
32    }
33    %mapB = rise.mapSeq #rise.nat<1024> #rise.array<1024, scalar<f32>> #rise.array<1024, scalar<f32>>
34    %result = rise.apply %mapB, %f2, %B
35    rise.return %result : !rise.array<1024, array<1024, scalar<f32>>>
36  }
37  %mapA = rise.mapSeq #rise.nat<1024> #rise.array<1024, scalar<f32>> #rise.array<1024, scalar<f32>>
38  %result = rise.apply %mapA, %f1, %A
39  rise.out %out ← %result
40  return
41 }
```


Patterns: zip

```
6 %zip = rise.zip #rise.nat<1024> #rise.scalar<f32> #rise.scalar<f32>
```

Patterns: zip

```
6 %zip = rise.zip #rise.nat<1024> #rise.scalar<f32> #rise.scalar<f32>

// type: () → !rise.fun<array<1024, scalar<f32>> →
                                fun<array<1024, scalar<f32>> →
                                    array<1024, tuple<scalar<f32>, scalar<f32>>>>>>

// type: () → 
```

- Each **RISE** pattern is implemented as an MLIR operation
- Operations customized with attributes specifying information for the type
- Patterns encoded as operations have a **RISE** function type

Function application

4

5

6

7

```
%zip = rise.zip #rise.nat<1024> #rise.scalar<f32> #rise.scalar<f32>
```

Function application

```
4 %f1 = rise.lambda (%arow : !rise.array<1024, scalar<f32>>) →  
    array<1024, scalar<f32>> {  
5     %f2 = rise.lambda (%bcol : !rise.array<1024, scalar<f32>>) →  
        array<1024, scalar<f32>> {  
6         %zip = rise.zip #rise.nat<1024> #rise.scalar<f32> #rise.scalar<f32>  
7         %zipped = rise.apply %zip, %arow, %bcol
```

- **rise.apply** models function application:
expects an SSA value with a **Rise function type** (%zip)
and arguments to the function (%arow, %bcol)

Function application

```
4 %f1 = rise.lambda (%arow : !rise.array<1024, scalar<f32>>) →  
    array<1024, scalar<f32>> {  
5   %f2 = rise.lambda (%bcol : !rise.array<1024, scalar<f32>>) →  
    array<1024, scalar<f32>> {  
6     %zip = rise.zip #rise.nat<1024> #rise.scalar<f32> #rise.scalar<f32>  
7     %zipped = rise.apply %zip, %arow, %bcol  
  
//type: (%zip.type, array<1024, scalar<f32>>, array<1024, scalar<f32>>) →  
    array<1024, tuple<scalar<f32>, scalar<f32>>>
```

- **rise.apply** models function application:
expects an SSA value with a **Rise function type** (%zip)
and arguments to the function (%arow, %bcol)

Function abstraction, aka λ -expressions

```
1 func @mm(%out:memref<1024×1024xf32>, %inA:memref<1024×1024xf32>, %inB:memref<1024×1024xf32>) {
2   %A = rise.in %inA : !rise.array<1024, array<1024, scalar<f32>>>
3   %B = rise.in %inB : !rise.array<1024, array<1024, scalar<f32>>>
4   %f1 = rise.lambda (%arow : !rise.array<1024, scalar<f32>>) → !rise.array<1024, scalar<f32>> {
5     %f2 = rise.lambda (%bcol : !rise.array<1024, scalar<f32>>) → !rise.array<1024, scalar<f32>> {
6       %zip = rise.zip #rise.nat<1024> #rise.scalar<f32> #rise.scalar<f32>
7       %zipped = rise.apply %zip, %arow, %bcol
8       %f = rise.lambda (%tuple : !rise.tuple<scalar<f32>, scalar<f32>>) → !rise.scalar<f32> {
9         %fstFun = rise.fst #rise.scalar<f32> #rise.scalar<f32>
10        %sndFun = rise.snd #rise.scalar<f32> #rise.scalar<f32>
11        %fst = rise.apply %fstFun, %tuple
12        %snd = rise.apply %sndFun, %tuple
13        %result = rise.embed(%fst, %snd) {
14          %res = mulf %fst, %snd : f32
15          return %res : f32
16        } : !rise.scalar<f32>
17        rise.return %result : !rise.scalar<f32>
18      }
19      %map = rise.mapSeq #rise.nat<1024> #rise.tuple<scalar<f32>, scalar<f32>> #rise.scalar<f32>
20      %multipliedArray = rise.apply %map, %f, %zipped
21      %add = rise.lambda (%a : !rise.scalar<f32>, %b : !rise.scalar<f32>) → !rise.scalar<f32>> {
22        %result = rise.embed(%a, %b) {
23          %res = addf %a, %b : f32
24          return %res : f32
25        } : !rise.scalar<f32>
26        rise.return %result : !rise.scalar<f32>
27      }
28      %init = rise.literal #rise.lit<0.0>
29      %reduce = rise.reduceSeq #rise.nat<1024> #rise.scalar<f32> #rise.scalar<f32>
30      %result = rise.apply %reduce, %add, %init, %multipliedArray
31      rise.return %result : !rise.scalar<f32>
32    }
33    %mapB = rise.mapSeq #rise.nat<1024> #rise.array<1024, scalar<f32>> #rise.array<1024, scalar<f32>>
34    %result = rise.apply %mapB, %f2, %B
35    rise.return %result : !rise.array<1024, array<1024, scalar<f32>>>
36  }
37  %mapA = rise.mapSeq #rise.nat<1024> #rise.array<1024, scalar<f32>> #rise.array<1024, scalar<f32>>
38  %result = rise.apply %mapA, %f1, %A
39  rise.out %out ← %result
40  return
41 }
```

Function abstraction, aka λ -expressions

```
1 func @mm(%out:memref<1024×1024xf32>, %inA:memref<1024×1024xf32>, %inB:memref<1024×1024xf32>) {
2   %A = rise.in %inA : !rise.array<1024, array<1024, scalar<f32>>>
3   %B = rise.in %inB : !rise.array<1024, array<1024, scalar<f32>>>
4   %f1 = rise.lambda (%arow : !rise.array<1024, scalar<f32>>) → !rise.array<1024, scalar<f32>> {
5     %f2 = rise.lambda (%bcol : !rise.array<1024, scalar<f32>>) → !rise.array<1024, scalar<f32>> {
6       %zip = rise.zip #rise.nat<1024> #rise.scalar<f32> #rise.scalar<f32>
7       %zipped = rise.apply %zip, %arow, %bcol
8       %f = rise.lambda (%tuple : !rise.tuple<scalar<f32>, scalar<f32>>) → !rise.scalar<f32> {
9         %fstFun = rise.fst #rise.scalar<f32> #rise.scalar<f32>
10        %sndFun = rise.snd #rise.scalar<f32> #rise.scalar<f32>
11        %fst = rise.apply %fstFun, %tuple
12        %snd = rise.apply %sndFun, %tuple
13        %result = rise.embed(%fst, %snd) {
14          %res = mulf %fst, %snd : f32
15          return %res : f32
16        } : !rise.scalar<f32>
17        rise.return %result : !rise.scalar<f32>
18      }
19      %map = rise.mapSeq #rise.nat<1024> #rise.tuple<scalar<f32>, scalar<f32>> #rise.scalar<f32>
20      %multipliedArray = rise.apply %map, %f, %zipped
21      %add = rise.lambda (%a : !rise.scalar<f32>, %b : !rise.scalar<f32>) → !rise.scalar<f32>> {
22        %result = rise.embed(%a, %b) {
23          %res = addf %a, %b : f32
24          return %res : f32
25        } : !rise.scalar<f32>
26        rise.return %result : !rise.scalar<f32>
27      }
28      %init = rise.lit<0.0>
29      %reduce = rise.reduceSeq #rise.nat<1024> #rise.scalar<f32> #rise.scalar<f32>
30      %result = rise.apply %reduce, %add, %init, %multipliedArray
31      rise.return %result : !rise.scalar<f32>
32    }
33    %mapB = rise.mapSeq #rise.nat<1024> #rise.array<1024, scalar<f32>> #rise.array<1024, scalar<f32>>
34    %result = rise.apply %mapB, %f2, %B
35    rise.return %result : !rise.array<1024, array<1024, scalar<f32>>>
36  }
37  %mapA = rise.mapSeq #rise.nat<1024> #rise.array<1024, scalar<f32>> #rise.array<1024, scalar<f32>>
38  %result = rise.apply %mapA, %f1, %A
39  rise.out %out ← %result
40  return
41 }
```

Function abstraction, aka λ -expressions

```
21 %add = rise.lambda (%a : !rise.scalar<f32>, %b : !rise.scalar<f32>) → !rise.scalar<f32> {  
22     %result = rise.embed(%a, %b) {  
23         %res = addf %a, %b : f32  
24         return %res  
25     } : !rise.scalar<f32>  
26     rise.return %result : !rise.scalar<f32>  
27 }
```

- `rise.lambda` has an MLIR region of exactly one block
- Arbitrary number of arguments and a result
- `rise.lambda` associates the region with a **RISE function type**

Function abstraction, aka λ -expressions

```
21 %add = rise.lambda (%a : !rise.scalar<f32>, %b : !rise.scalar<f32>) → !rise.scalar<f32> {  
22     %result = rise.embed(%a, %b) {  
23         %res = addf %a, %b : f32  
24         return %res  
25     } : !rise.scalar<f32>  
26     rise.return %result : !rise.scalar<f32>  
27 } // type: !rise.fun<scalar<f32> → fun<scalar<f32> → scalar<f32>>>
```

- `rise.lambda` has an MLIR region of exactly one block
- Arbitrary number of arguments and a result
- `rise.lambda` associates the region with a **RISE function type**

Embedding of other dialects

```
21 %add = rise.lambda (%a : !rise.scalar<f32>, %b : !rise.scalar<f32>) → !rise.scalar<f32> {  
22     %result = rise.embed(%a, %b) {  
23         %res = addf %a, %b : f32  
24         return %res  
25     } : !rise.scalar<f32>  
26     rise.return %result : !rise.scalar<f32>  
27 }
```

- `rise.embed` unwraps the operands of type `!rise.scalar<t>` and exposes the values of the wrapped types `t` inside the region
- The region may contain operations from `arbitrary` MLIR dialects
- `rise.embed` returns the result value of `!rise.scalar` type

Embedding of other dialects

```
21 %add = rise.lambda (%a : !rise.scalar<f32>, %b : !rise.scalar<f32>) → !rise.scalar<f32> {  
    // %a, %b : !rise.scalar<f32>  
22     %result = rise.embed(%a, %b) {  
        // %a, %b : f32  
23         %res = addf %a, %b : f32  
24         return %res  
25     } : !rise.scalar<f32>  
26     rise.return %result : !rise.scalar<f32>  
27 }
```

- `rise.embed` unwraps the operands of type `!rise.scalar<t>` and exposes the values of the wrapped types `t` inside the region
- The region may contain operations from `arbitrary` MLIR dialects
- `rise.embed` returns the result value of `!rise.scalar` type

RISE dialect by example: Matrix Multiplication

```
1 func @mm(%out:memref<1024×1024xf32>, %inA:memref<1024×1024xf32>, %inB:memref<1024×1024xf32>) {
2   %A = rise.in %inA : !rise.array<1024, array<1024, scalar<f32>>>
3   %B = rise.in %inB : !rise.array<1024, array<1024, scalar<f32>>>
4   %f1 = rise.lambda (%arow : !rise.array<1024, scalar<f32>>) → !rise.array<1024, scalar<f32>> {
5     %f2 = rise.lambda (%bcol : !rise.array<1024, scalar<f32>>) → !rise.array<1024, scalar<f32>> {
6       %zip = rise.zip #rise.nat<1024> #rise.scalar<f32> #rise.scalar<f32>
7       %zipped = rise.apply %zip, %arow, %bcol
8       %f = rise.lambda (%tuple : !rise.tuple<scalar<f32>, scalar<f32>>) → !rise.scalar<f32> {
9         %fstFun = rise.fst #rise.scalar<f32> #rise.scalar<f32>
10        %sndFun = rise.snd #rise.scalar<f32> #rise.scalar<f32>
11        %fst = rise.apply %fstFun, %tuple
12        %snd = rise.apply %sndFun, %tuple
13        %result = rise.embed(%fst, %snd) {
14          %res = mulf %fst, %snd : f32
15          return %res : f32
16        } : !rise.scalar<f32>
17        rise.return %result : !rise.scalar<f32>
18      }
19      %map = rise.mapSeq #rise.nat<1024> #rise.tuple<scalar<f32>, scalar<f32>> #rise.scalar<f32>
20      %multipliedArray = rise.apply %map, %f, %zipped
21      %add = rise.lambda (%a : !rise.scalar<f32>, %b : !rise.scalar<f32>) → !rise.scalar<f32>> {
22        %result = rise.embed(%a, %b) {
23          %res = addf %a, %b : f32
24          return %res : f32
25        } : !rise.scalar<f32>
26        rise.return %result : !rise.scalar<f32>
27      }
28      %init = rise.literal #rise.lit<0.0>
29      %reduce = rise.reduceSeq #rise.nat<1024> #rise.scalar<f32> #rise.scalar<f32>
30      %result = rise.apply %reduce, %add, %init, %multipliedArray
31      rise.return %result : !rise.scalar<f32>
32    }
33    %mapB = rise.mapSeq #rise.nat<1024> #rise.array<1024, scalar<f32>> #rise.array<1024, scalar<f32>>
34    %result = rise.apply %mapB, %f2, %B
35    rise.return %result : !rise.array<1024, array<1024, scalar<f32>>>
36  }
37  %mapA = rise.mapSeq #rise.nat<1024> #rise.array<1024, scalar<f32>> #rise.array<1024, scalar<f32>>
38  %result = rise.apply %mapA, %f1, %A
39  rise.out %out ← %result
40  return
41 }
```


RISE dialect by example: Matrix Multiplication

```

1 func @mm(%out:memref<1024×1024xf32>, %inA:memref<1024×1024xf32>, %inB:memref<1024×1024xf32>) {
2   %A = rise.in %inA : !rise.array<1024, array<1024, scalar<f32>>>
3   %B = rise.in %inB : !rise.array<1024, array<1024, scalar<f32>>>
4   %f1 = rise.lambda (%arow : !rise.array<1024, scalar<f32>>) → !rise.array<1024, scalar<f32>> {
5     %f2 = rise.lambda (%bcol : !rise.array<1024, scalar<f32>>) → !rise.array<1024, scalar<f32>> {
6       %zip = rise.zip #rise.nat<1024> #rise.scalar<f32> #rise.scalar<f32>
7       %zipped = rise.apply %zip, %arow, %bcol
8       %f = rise.lambda (%tuple : !rise.tuple<scalar<f32>, scalar<f32>>) → !rise.scalar<f32> {
9         %fstFun = rise.fst #rise.scalar<f32> #rise.scalar<f32>
10        %sndFun = rise.snd #rise.scalar<f32> #rise.scalar<f32>
11        %fst = rise.apply %fstFun, %tuple
12        %snd = rise.apply %sndFun, %tuple
13        %result = rise.embed(%fst, %snd) {
14          %res = mulf %fst, %snd : f32
15          return %res : f32
16        } : !rise.scalar<f32>
17        rise.return %result : !rise.scalar<f32>
18      }
19      %map = rise.mapSeq #rise.nat<1024> #rise.tuple<scalar<f32>, scalar<f32>> #rise.scalar<f32>
20      %multipliedArray = rise.apply %map, %f, %zipped
21      %add = rise.lambda (%a : !rise.scalar<f32>, %b : !rise.scalar<f32>) → !rise.scalar<f32> {
22        %result = rise.embed(%a, %b) {
23          %res = addf %a, %b : f32
24          return %res : f32
25        } : !rise.scalar<f32>
26        rise.return %result : !rise.scalar<f32>
27      }
28      %init = rise.literal #rise.lit<0.0>
29      %reduce = rise.reduceSeq #rise.nat<1024> #rise.scalar<f32> #rise.scalar<f32>
30      %result = rise.apply %reduce, %add, %init, %multipliedArray
31      rise.return %result : !rise.scalar<f32>
32    }
33    %mapB = rise.mapSeq #rise.nat<1024> #rise.array<1024, scalar<f32>> #rise.array<1024, scalar<f32>>
34    %result = rise.apply %mapB, %f2, %B
35    rise.return %result : !rise.array<1024, array<1024, scalar<f32>>>
36  }
37  %mapA = rise.mapSeq #rise.nat<1024> #rise.array<1024, scalar<f32>> #rise.array<1024, scalar<f32>>
38  %result = rise.apply %mapA, %f1, %A
39  rise.out %out ← %result
40  return
41 }

```

mapSeq (
mapSeq (



R|SE dialect - modelling λ -calculus + patterns

Core λ -calculus: `rise.lambda`, `rise.apply`, *`rise.let`*,
`!rise.array`, *`!rise.tuple`*, *`!rise.fun`*

Patterns: `rise.zip`, `rise.mapSeq`, `rise.reduceSeq`, *`rise.mapPar`*, *`rise.reducePar`*,
`rise.literal` `rise.tuple`, `rise.fst`, `rise.snd`,
`rise.split`, *`rise.join`*, *`rise.transpose`*, ...

Interoperability: `rise.embed`, `rise.in`, `rise.out`

-> direct correspondence of MLIR dialect with original functional representation

Lowering of the RISE dialect

Lowering of patterns is not context free, e.g. zip influences the code generation of the next pattern

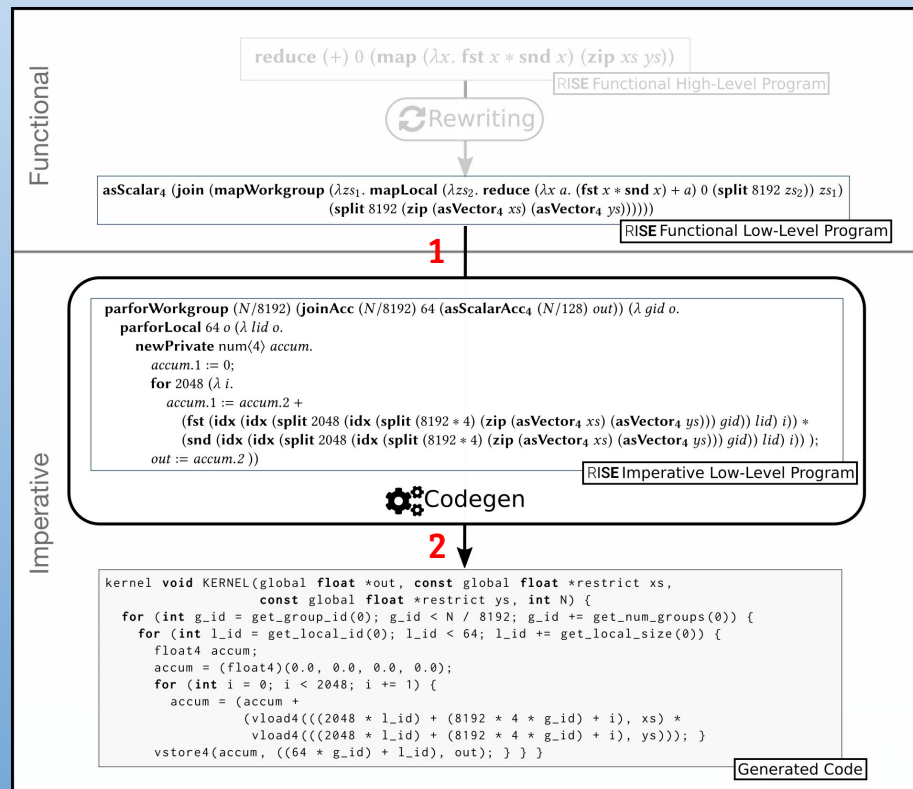
Lowering is divided into two steps:

1. Lowering functional to imperative
2. Lowering to final target code

Lowering described formally in:

Strategy Preserving Compilation for Parallel Functional Code
by Robert Atkey, Michel Steuwer, Sam Lindley, and Christophe Dubach

<https://arxiv.org/abs/1710.08332>



Lowering of the RISE dialect

Lowering of patterns is not context free, e.g. zip influences the code generation of the next pattern

Lowering is divided into two steps:

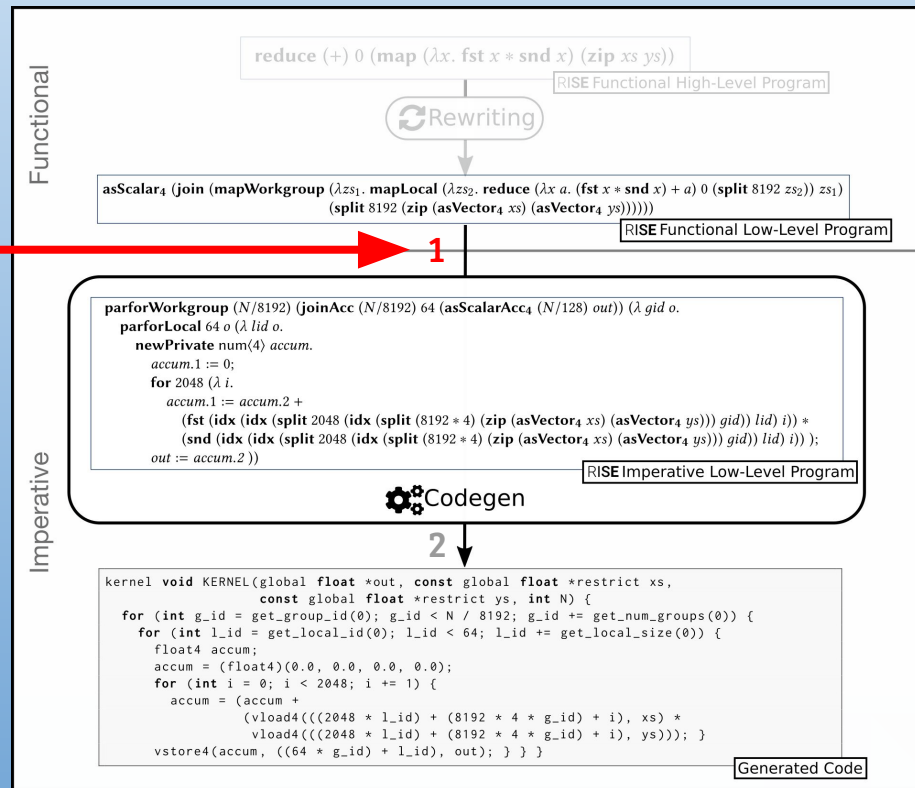
1. Lowering functional to imperative

2. Lowering to final target code

Lowering described formally in:

Strategy Preserving Compilation for Parallel Functional Code
by Robert Atkey, Michel Steuwer, Sam Lindley, and Christophe Dubach

<https://arxiv.org/abs/1710.08332>



Lowering to rise.codegen: dot

```
1 func @dot(%out:memref<1xf32>, %a:memref<1024xf32>,
           %b:memref<1024xf32>) {
2   %lhs = rise.in %a : !rise.array<1024, scalar<f32>>
3   %rhs = rise.in %b : !rise.array<1024, scalar<f32>>
4
5   %zip = rise.zip #rise.nat<1024> #rise.scalar<f32>
           #rise.scalar<f32>
6   %zipped = rise.apply %zip, %lhs, %rhs
7   %f = rise.lambda (%tuple, %acc) → !rise.scalar<f32> {
8     %fstFun = rise.fst #rise.scalar<f32> #rise.scalar<f32>
9     %sndFun = rise.snd #rise.scalar<f32> #rise.scalar<f32>
10    %fst = rise.apply %fstFun, %tuple
11    %snd = rise.apply %sndFun, %tuple
12    %result = rise.embed(%fst, %snd, %acc) {
13      %product = mulf %fst, %snd : f32
14      %res = addf %product, %acc : f32
15      return res : f32
16    } : !rise.scalar<f32>
17    rise.return %result : !rise.scalar<f32>
18  }
19  %init = rise.literal #rise.lit<0.0>
20  %reduce = rise.reduceSeq #rise.nat<1024>
           #rise.tuple<scalar<f32>, scalar<f32>> #rise.scalar<f32>
21  %result = rise.apply %reduce, %f, %init, %zipped
22  rise.out %out ← %result
23  return
24 }
```

RISE

lowering

```
1 func @dot(%out:memref<1xf32>,%a:memref<1024xf32>,%b:memref<1024xf32>) {
2   %cst0 = constant 0.000000e+00 : f32
3   %c0 = constant 0 : index
4   %c1024 = constant 1024 : index
5   %c1 = constant 1 : index
6   %0 = rise.codegen.zip(%a, %b)
           : (memref<1024xf32>, memref<1024xf32>) → memref<1024xf32>
7   %1 = rise.codegen.idx(%out, %c0) : (memref<1xf32>, index) → f32
8   rise.codegen.assign(%cst0, %1) : (f32, f32) → ()
9   loop.for %i = %c0 to %c1024 step %c1 {
10    %2 = rise.codegen.idx(%0, %i) : (memref<1024xf32>, index) → f32
11    %3 = rise.codegen.idx(%out, %c0) : (memref<1xf32>, index) → f32
12    %4 = rise.codegen.fst(%2) : (f32) → f32
13    %5 = rise.codegen.snd(%2) : (f32) → f32
14    %6 = mulf %4, %5 : f32
15    %7 = addf %6, %3 : f32
16    rise.codegen.assign(%7, %3) : (f32, f32) → ()
17  }
18  return
19 }
```

Loop + RISE Intermediate

Lowering to rise.codegen: dot

```
1 func @dot(%out:memref<1xf32>, %a:memref<1024xf32>,
           %b:memref<1024xf32>) {
2   %lhs = rise.in %a : !rise.array<1024, scalar<f32>>
3   %rhs = rise.in %b : !rise.array<1024, scalar<f32>>
4
5   %zip = rise.zip #rise.nat<1024> #rise.scalar<f32>
           #rise.scalar<f32>
6   %zipped = rise.apply %zip, %lhs, %rhs
7   %f = rise.lambda (%tuple, %acc) → !rise.scalar<f32> {
8     %fstFun = rise.fst #rise.scalar<f32> #rise.scalar<f32>
9     %sndFun = rise.snd #rise.scalar<f32> #rise.scalar<f32>
10    %fst = rise.apply %fstFun, %tuple
11    %snd = rise.apply %sndFun, %tuple
12    %result = rise.embed(%fst, %snd, %acc) {
13      %product = mulf %fst, %snd : f32
14      %res = addf %product, %acc : f32
15      return res : f32
16    } : !rise.scalar<f32>
17    rise.return %result : !rise.scalar<f32>
18  }
19  %init = rise.literal #rise.lit<0.0>
20  %reduce = rise.reduceSeq #rise.nat<1024>
           #rise.tuple<scalar<f32>, scalar<f32>> #rise.scalar<f32>
21  %result = rise.apply %reduce, %f, %init, %zipped
22  rise.out %out ← %result
23  return
24 }
```

RISE

lowering

```
1 func @dot(%out:memref<1xf32>,%a:memref<1024xf32>,%b:memref<1024xf32>) {
2   %cst0 = constant 0.000000e+00 : f32
3   %c0 = constant 0 : index
4   %c1024 = constant 1024 : index
5   %c1 = constant 1 : index
6   %0 = rise.codegen.zip(%a, %b)
           : (memref<1024xf32>, memref<1024xf32>) → memref<1024xf32>
7   %1 = rise.codegen.idx(%out, %c0) : (memref<1xf32>, index) → f32
8   rise.codegen.assign(%cst0, %1) : (f32, f32) → ()
9   loop.for %i = %c0 to %c1024 step %c1 {
10    %2 = rise.codegen.idx(%0, %i) : (memref<1024xf32>, index) → f32
11    %3 = rise.codegen.idx(%out, %c0) : (memref<1xf32>, index) → f32
12    %4 = rise.codegen.fst(%2) : (f32) → f32
13    %5 = rise.codegen.snd(%2) : (f32) → f32
14    %6 = mulf %4, %5 : f32
15    %7 = addf %6, %3 : f32
16    rise.codegen.assign(%7, %3) : (f32, f32) → ()
17  }
18  return
19 }
```

Loop + RISE Intermediate

- Start translation from rise.out

Lowering to rise.codegen: dot

```
1 func @dot(%out:memref<1xf32>, %a:memref<1024xf32>,
           %b:memref<1024xf32>) {
2   %lhs = rise.in %a : !rise.array<1024, scalar<f32>>
3   %rhs = rise.in %b : !rise.array<1024, scalar<f32>>
4
5   %zip = rise.zip #rise.nat<1024> #rise.scalar<f32>
           #rise.scalar<f32>
6   %zipped = rise.apply %zip, %lhs, %rhs
7   %f = rise.lambda (%tuple, %acc) → !rise.scalar<f32> {
8     %fstFun = rise.fst #rise.scalar<f32> #rise.scalar<f32>
9     %sndFun = rise.snd #rise.scalar<f32> #rise.scalar<f32>
10    %fst = rise.apply %fstFun, %tuple
11    %snd = rise.apply %sndFun, %tuple
12    %result = rise.embed(%fst, %snd, %acc) {
13      %product = mulf %fst, %snd : f32
14      %res = addf %product, %acc : f32
15      return res : f32
16    } : !rise.scalar<f32>
17    rise.return %result : !rise.scalar<f32>
18  }
19  %init = rise.literal #rise.lit<0.0>
20  %reduce = rise.reduceSeq #rise.nat<1024>
           #rise.tuple<scalar<f32>, scalar<f32>> #rise.scalar<f32>
21  %result = rise.apply %reduce, %f, %init, %zipped
22  rise.out %out ← %result
23  return
24 }
```

RISE

lowering

```
1 func @dot(%out:memref<1xf32>,%a:memref<1024xf32>,%b:memref<1024xf32>) {
2   %cst0 = constant 0.000000e+00 : f32
3   %c0 = constant 0 : index
4   %c1024 = constant 1024 : index
5   %c1 = constant 1 : index
6   %0 = rise.codegen.zip(%a, %b)
           : (memref<1024xf32>, memref<1024xf32>) → memref<1024xf32>
7   %1 = rise.codegen.idx(%out, %c0) : (memref<1xf32>, index) → f32
8   rise.codegen.assign(%cst0, %1) : (f32, f32) → ()
9   loop.for %i = %c0 to %c1024 step %c1 {
10    %2 = rise.codegen.idx(%0, %i) : (memref<1024xf32>, index) → f32
11    %3 = rise.codegen.idx(%out, %c0) : (memref<1xf32>, index) → f32
12    %4 = rise.codegen.fst(%2) : (f32) → f32
13    %5 = rise.codegen.snd(%2) : (f32) → f32
14    %6 = mulf %4, %5 : f32
15    %7 = addf %6, %3 : f32
16    rise.codegen.assign(%7, %3) : (f32, f32) → ()
17  }
18  return
19 }
```

Loop + RISE Intermediate

- Start translation from `rise.out`
- Traverse program back to front and lower patterns as specified by their arguments

Lowering to rise.codegen: dot

```

1 func @dot(%out:memref<1xf32>, %a:memref<1024xf32>,
           %b:memref<1024xf32>) {
2   %lhs = rise.in %a : !rise.array<1024, scalar<f32>>
3   %rhs = rise.in %b : !rise.array<1024, scalar<f32>>
4
5   %zip = rise.zip #rise.nat<1024> #rise.scalar<f32>
           #rise.scalar<f32>
6   %zipped = rise.apply %zip, %lhs, %rhs
7   %f = rise.lambda (%tuple, %acc) → !rise.scalar<f32> {
8     %fstFun = rise.fst #rise.scalar<f32> #rise.scalar<f32>
9     %sndFun = rise.snd #rise.scalar<f32> #rise.scalar<f32>
10    %fst = rise.apply %fstFun, %tuple
11    %snd = rise.apply %sndFun, %tuple
12    %result = rise.embed(%fst, %snd, %acc) {
13      %product = mulf %fst, %snd : f32
14      %res = addf %product, %acc : f32
15      return res : f32
16    } : !rise.scalar<f32>
17    rise.return %result : !rise.scalar<f32>
18  }
19  %init = rise.literal #rise.lit<0.0>
20  %reduce = rise.reduceSeq #rise.nat<1024>
           #rise.tuple<scalar<f32>, scalar<f32>> #rise.scalar<f32>
21  %result = rise.apply %reduce, %f, %init, %zipped
22  rise.out %out ← %result
23  return
24 }

```

RISE

lowering
reduceSeq

```

1 func @dot(%out:memref<1xf32>,%a:memref<1024xf32>,%b:memref<1024xf32>) {
2   %cst0 = constant 0.000000e+00 : f32
3   %c0 = constant 0 : index
4   %c1024 = constant 1024 : index
5   %c1 = constant 1 : index
6   %0 = rise.codegen.zip(%a, %b)
           : (memref<1024xf32>, memref<1024xf32>) → memref<1024xf32>
7   %1 = rise.codegen.idx(%out, %c0) : (memref<1xf32>, index) → f32
8   rise.codegen.assign(%cst0, %1) : (f32, f32) → ()
9   loop.for %i = %c0 to %c1024 step %c1 {
10    %2 = rise.codegen.idx(%0, %i) : (memref<1024xf32>, index) → f32
11    %3 = rise.codegen.idx(%out, %c0) : (memref<1xf32>, index) → f32
12    %4 = rise.codegen.fst(%2) : (f32) → f32
13    %5 = rise.codegen.snd(%2) : (f32) → f32
14    %6 = mulf %4, %5 : f32
15    %7 = addf %6, %3 : f32
16    rise.codegen.assign(%7, %3) : (f32, f32) → ()
17  }
18  return
19 }

```

Loop + RISE Intermediate

$$\begin{aligned}
 out :=_{\delta_2} \text{reduceSeq}(N, \delta_1, \delta_2, F, I, E) &= C(I)_{\delta_2}(\lambda \text{init}. \\
 &C(E)_{N, \delta_1}(\lambda x. \\
 &\quad out :=_{\delta_1} \text{init}; \\
 &\quad \text{for } N (\lambda i. out :=_{\delta_2} F(x[i], out));))
 \end{aligned}$$

- Start translation from `rise.out`
- Traverse program back to front and lower patterns as specified by their arguments

Lowering to rise.codegen: dot

```

1 func @dot(%out:memref<1xf32>, %a:memref<1024xf32>,
           %b:memref<1024xf32>) {
2   %lhs = rise.in %a : !rise.array<1024, scalar<f32>>
3   %rhs = rise.in %b : !rise.array<1024, scalar<f32>>
4
5   %zip = rise.zip #rise.nat<1024> #rise.scalar<f32>
           #rise.scalar<f32>
6   %zipped = rise.apply %zip, %lhs, %rhs
7   %f = rise.lambda (%tuple, %acc) → !rise.scalar<f32> {
8     %fstFun = rise.fst #rise.scalar<f32> #rise.scalar<f32>
9     %sndFun = rise.snd #rise.scalar<f32> #rise.scalar<f32>
10    %fst = rise.apply %fstFun, %tuple
11    %snd = rise.apply %sndFun, %tuple
12    %result = rise.embed(%fst, %snd, %acc) {
13      %product = mulf %fst, %snd : f32
14      %res = addf %product, %acc : f32
15      return res : f32
16    } : !rise.scalar<f32>
17    rise.return %result : !rise.scalar<f32>
18  }
19  %init = rise.literal #rise.lit<0.0>
20  %reduce = rise.reduceSeq #rise.nat<1024>
           #rise.tuple<scalar<f32>, scalar<f32>> #rise.scalar<f32>
21  %result = rise.apply %reduce, %f, %init, %zipped
22  rise.out %out ← %result
23  return
24 }

```

RISE

lowering
reduceSeq

```

1 func @dot(%out:memref<1xf32>,%a:memref<1024xf32>,%b:memref<1024xf32>) {
2   %cst0 = constant 0.000000e+00 : f32
3   %c0 = constant 0 : index
4   %c1024 = constant 1024 : index
5   %c1 = constant 1 : index
6   %0 = rise.codegen.zip(%a, %b)
           : (memref<1024xf32>, memref<1024xf32>) → memref<1024xf32>
7   %1 = rise.codegen.idx(%out, %c0) : (memref<1xf32>, index) → f32
8   rise.codegen.assign(%cst0, %1) : (f32, f32) → ()
9   loop.for %i = %c0 to %c1024 step %c1 {
10    %2 = rise.codegen.idx(%0, %i) : (memref<1024xf32>, index) → f32
11    %3 = rise.codegen.idx(%out, %c0) : (memref<1xf32>, index) → f32
12    %4 = rise.codegen.fst(%2) : (f32) → f32
13    %5 = rise.codegen.snd(%2) : (f32) → f32
14    %6 = mulf %4, %5 : f32
15    %7 = addf %6, %3 : f32
16    rise.codegen.assign(%7, %3) : (f32, f32) → ()
17  }
18  return
19 }

```

Loop + RISE Intermediate

$$\begin{aligned}
 out :=_{\delta_2} \text{reduceSeq}(N, \delta_1, \delta_2, F, I, E) &= C(I)_{\delta_2}(\lambda \text{init}. \\
 &C(E)_{N, \delta_1}(\lambda x. \\
 &\quad out :=_{\delta_1} \text{init}; \\
 &\quad \text{for } N(\lambda i. out :=_{\delta_2} F(x[i], out));))
 \end{aligned}$$

- Start translation from `rise.out`
- Traverse program back to front and lower patterns as specified by their arguments

Lowering to rise.codegen: dot

```

1 func @dot(%out:memref<1xf32>, %a:memref<1024xf32>,
           %b:memref<1024xf32>) {
2   %lhs = rise.in %a : !rise.array<1024, scalar<f32>>
3   %rhs = rise.in %b : !rise.array<1024, scalar<f32>>
4
5   %zip = rise.zip #rise.nat<1024> #rise.scalar<f32>
           #rise.scalar<f32>
6   %zipped = rise.apply %zip, %lhs, %rhs
7   %f = rise.lambda (%tuple, %acc) → !rise.scalar<f32> {
8     %fstFun = rise.fst #rise.scalar<f32> #rise.scalar<f32>
9     %sndFun = rise.snd #rise.scalar<f32> #rise.scalar<f32>
10    %fst = rise.apply %fstFun, %tuple
11    %snd = rise.apply %sndFun, %tuple
12    %result = rise.embed(%fst, %snd, %acc) {
13      %product = mulf %fst, %snd : f32
14      %res = addf %product, %acc : f32
15      return res : f32
16    } : !rise.scalar<f32>
17    rise.return %result : !rise.scalar<f32>
18  }
19  %init = rise.literal #rise.lit<0.0>
20  %reduce = rise.reduceSeq #rise.nat<1024>
           #rise.tuple<scalar<f32>, scalar<f32>> #rise.scalar<f32>
21  %result = rise.apply %reduce, %f, %init, %zipped
22  rise.out %out ← %result
23  return
24 }

```

RISE

lowering
literal

```

1 func @dot(%out:memref<1xf32>,%a:memref<1024xf32>,%b:memref<1024xf32>) {
2   %cst0 = constant 0.000000e+00 : f32
3   %c0 = constant 0 : index
4   %c1024 = constant 1024 : index
5   %c1 = constant 1 : index
6   %0 = rise.codegen.zip(%a, %b)
           : (memref<1024xf32>, memref<1024xf32>) → memref<1024xf32>
7   %1 = rise.codegen.idx(%out, %c0) : (memref<1xf32>, index) → f32
8   rise.codegen.assign(%cst0, %1) : (f32, f32) → ()
9   loop.for %i = %c0 to %c1024 step %c1 {
10    %2 = rise.codegen.idx(%0, %i) : (memref<1024xf32>, index) → f32
11    %3 = rise.codegen.idx(%out, %c0) : (memref<1xf32>, index) → f32
12    %4 = rise.codegen.fst(%2) : (f32) → f32
13    %5 = rise.codegen.snd(%2) : (f32) → f32
14    %6 = mulf %4, %5 : f32
15    %7 = addf %6, %3 : f32
16    rise.codegen.assign(%7, %3) : (f32, f32) → ()
17  }
18  return
19 }

```

Loop + RISE Intermediate

$$out :=_{\delta_2} \text{reduceSeq}(N, \delta_1, \delta_2, F, \boxed{I}, E) = \boxed{C(\boxed{I})_{\delta_2}(\lambda \text{init}.}$$

$$\boxed{C(E)_{N.\delta_1}(\lambda x.}$$

$$\boxed{out :=_{\delta_1} \text{init};}$$

$$\text{for } N (\lambda i. out :=_{\delta_2} F(x[i], out));))$$

- Start translation from `rise.out`
- Traverse program back to front and lower patterns as specified by their arguments

Lowering to rise.codegen: dot

```

1 func @dot(%out:memref<1xf32>, %a:memref<1024xf32>,
           %b:memref<1024xf32>) {
2   %lhs = rise.in %a : !rise.array<1024, scalar<f32>>
3   %rhs = rise.in %b : !rise.array<1024, scalar<f32>>
4
5   %zip = rise.zip #rise.nat<1024> #rise.scalar<f32>
        #rise.scalar<f32>
6   %zipped = rise.apply %zip, %lhs, %rhs
7   %f = rise.lambda (%tuple, %acc) → !rise.scalar<f32> {
8     %fstFun = rise.fst #rise.scalar<f32> #rise.scalar<f32>
9     %sndFun = rise.snd #rise.scalar<f32> #rise.scalar<f32>
10    %fst = rise.apply %fstFun, %tuple
11    %snd = rise.apply %sndFun, %tuple
12    %result = rise.embed(%fst, %snd, %acc) {
13      %product = mulf %fst, %snd : f32
14      %res = addf %product, %acc : f32
15      return res : f32
16    } : !rise.scalar<f32>
17    rise.return %result : !rise.scalar<f32>
18  }
19  %init = rise.literal #rise.lit<0.0>
20  %reduce = rise.reduceSeq #rise.nat<1024>
        #rise.tuple<scalar<f32>, scalar<f32>> #rise.scalar<f32>
21  %result = rise.apply %reduce, %f, %init, %zipped
22  rise.out %out ← %result
23  return
24 }

```

RISE

lowering

zip

```

1 func @dot(%out:memref<1xf32>,%a:memref<1024xf32>,%b:memref<1024xf32>) {
2   %cst0 = constant 0.000000e+00 : f32
3   %c0 = constant 0 : index
4   %c1024 = constant 1024 : index
5   %c1 = constant 1 : index
6   %0 = rise.codegen.zip(%a, %b)
        : (memref<1024xf32>, memref<1024xf32>) → memref<1024xf32>
7   %1 = rise.codegen.idx(%out, %c0) : (memref<1xf32>, index) → f32
8   rise.codegen.assign(%cst0, %1) : (f32, f32) → ()
9   loop.for %i = %c0 to %c1024 step %c1 {
10    %2 = rise.codegen.idx(%0, %i) : (memref<1024xf32>, index) → f32
11    %3 = rise.codegen.idx(%out, %c0) : (memref<1xf32>, index) → f32
12    %4 = rise.codegen.fst(%2) : (f32) → f32
13    %5 = rise.codegen.snd(%2) : (f32) → f32
14    %6 = mulf %4, %5 : f32
15    %7 = addf %6, %3 : f32
16    rise.codegen.assign(%7, %3) : (f32, f32) → ()
17  }
18  return
19 }

```

Loop + RISE Intermediate

$$\begin{aligned}
 out :=_{\delta_2} \text{reduceSeq}(N, \delta_1, \delta_2, F, I, \boxed{E}) &= C(I)_{\delta_2}(\lambda \text{init}. \\
 &\quad \boxed{C(E)}_{N, \delta_1}(\lambda x. \\
 &\quad out :=_{\delta_1} \text{init}; \\
 &\quad \text{for } N(\lambda i. out :=_{\delta_2} F(\boxed{x[i]}, out));))
 \end{aligned}$$

- Start translation from rise.out
- Traverse program back to front and lower patterns as specified by their arguments

Lowering to rise.codegen: dot

```

1 func @dot(%out:memref<1xf32>, %a:memref<1024xf32>,
           %b:memref<1024xf32>) {
2   %lhs = rise.in %a : !rise.array<1024, scalar<f32>>
3   %rhs = rise.in %b : !rise.array<1024, scalar<f32>>
4
5   %zip = rise.zip #rise.nat<1024> #rise.scalar<f32>
         #rise.scalar<f32>
6   %zipped = rise.apply %zip, %lhs, %rhs
7   %f = rise.lambda (%tuple, %acc) → !rise.scalar<f32> {
8     %fstFun = rise.fst #rise.scalar<f32> #rise.scalar<f32>
9     %sndFun = rise.snd #rise.scalar<f32> #rise.scalar<f32>
10    %fst = rise.apply %fstFun, %tuple
11    %snd = rise.apply %sndFun, %tuple
12    %result = rise.embed(%fst, %snd, %acc) {
13      %product = mulf %fst, %snd : f32
14      %res = addf %product, %acc : f32
15      return res : f32
16    } : !rise.scalar<f32>
17    rise.return %result : !rise.scalar<f32>
18  }
19  %init = rise.literal #rise.lit<0.0>
20  %reduce = rise.reduceSeq #rise.nat<1024>
         #rise.tuple<scalar<f32>, scalar<f32>> #rise.scalar<f32>
21  %result = rise.apply %reduce, %f, %init, %zipped
22  rise.out %out ← %result
23  return
24 }

```

RISE

lowering

zip

```

1 func @dot(%out:memref<1xf32>,%a:memref<1024xf32>,%b:memref<1024xf32>) {
2   %cst0 = constant 0.000000e+00 : f32
3   %c0 = constant 0 : index
4   %c1024 = constant 1024 : index
5   %c1 = constant 1 : index
6   %0 = rise.codegen.zip(%a, %b)
       : (memref<1024xf32>, memref<1024xf32>) → memref<1024xf32>
7   %1 = rise.codegen.idx(%out, %c0) : (memref<1xf32>, index) → f32
8   rise.codegen.assign(%cst0, %1) : (f32, f32) → ()
9   loop.for %i = %c0 to %c1024 step %c1 {
10    %2 = rise.codegen.idx(%0, %i) : (memref<1024xf32>, index) → f32
11    %3 = rise.codegen.idx(%out, %c0) : (memref<1xf32>, index) → f32
12    %4 = rise.codegen.fst(%2) : (f32) → f32
13    %5 = rise.codegen.snd(%2) : (f32) → f32
14    %6 = mulf %4, %5 : f32
15    %7 = addf %6, %3 : f32
16    rise.codegen.assign(%7, %3) : (f32, f32) → ()
17  }
18  return
19 }

```

Loop + RISE Intermediate

$$\begin{aligned}
 out :=_{\delta_2} \text{reduceSeq}(N, \delta_1, \delta_2, F, I, \boxed{E}) &= C(I)_{\delta_2}(\lambda \text{init}. \\
 &\quad \boxed{C(E)}_{N, \delta_1}(\lambda x. \\
 &\quad \text{out} :=_{\delta_1} \text{init}; \\
 &\quad \text{for } N(\lambda i. \text{out} :=_{\delta_2} F(\boxed{x[i]}, \text{out}));))
 \end{aligned}$$

- Start translation from `rise.out`
- Traverse program back to front and lower patterns as specified by their arguments

Lowering to rise.codegen: dot

```

1 func @dot(%out:memref<1xf32>, %a:memref<1024xf32>,
           %b:memref<1024xf32>) {
2   %lhs = rise.in %a : !rise.array<1024, scalar<f32>>
3   %rhs = rise.in %b : !rise.array<1024, scalar<f32>>
4
5   %zip = rise.zip #rise.nat<1024> #rise.scalar<f32>
         #rise.scalar<f32>
6   %zipped = rise.apply %zip, %lhs, %rhs
7   %f = rise.lambda (%tuple, %acc) → !rise.scalar<f32> {
8     %fstFun = rise.fst #rise.scalar<f32> #rise.scalar<f32>
9     %sndFun = rise.snd #rise.scalar<f32> #rise.scalar<f32>
10    %fst = rise.apply %fstFun, %tuple
11    %snd = rise.apply %sndFun, %tuple
12    %result = rise.embed(%fst, %snd, %acc) {
13      %product = mulf %fst, %snd : f32
14      %res = addf %product, %acc : f32
15      return res : f32
16    } : !rise.scalar<f32>
17    rise.return %result : !rise.scalar<f32>
18  }
19 %init = rise.literal #rise.lit<0.0>
20 %reduce = rise.reduceSeq #rise.nat<1024>
         #rise.tuple<scalar<f32>, scalar<f32>> #rise.scalar<f32>
21 %result = rise.apply %reduce, %f, %init, %zipped
22 rise.out %out ← %result
23 return
24 }

```

RISE

lowering

zip

```

1 func @dot(%out:memref<1xf32>,%a:memref<1024xf32>,%b:memref<1024xf32>) {
2   %cst0 = constant 0.000000e+00 : f32
3   %c0 = constant 0 : index
4   %c1024 = constant 1024 : index
5   %c1 = constant 1 : index
6   %0 = rise.codegen.zip(%a, %b)
         : (memref<1024xf32>, memref<1024xf32>) → memref<1024xf32>
7   %1 = rise.codegen.idx(%out, %c0) : (memref<1xf32>, index) → f32
8   rise.codegen.assign(%cst0, %1) : (f32, f32) → ()
9   loop.for %i = %c0 to %c1024 step %c1 {
10    %2 = rise.codegen.idx(%0, %i) : (memref<1024xf32>, index) → f32
11    %3 = rise.codegen.idx(%out, %c0) : (memref<1xf32>, index) → f32
12    %4 = rise.codegen.fst(%2) : (f32) → f32
13    %5 = rise.codegen.snd(%2) : (f32) → f32
14    %6 = mulf %4, %5 : f32
15    %7 = addf %6, %3 : f32
16    rise.codegen.assign(%7, %3) : (f32, f32) → ()
17  }
18  return
19 }

```

Loop + RISE Intermediate

$$\begin{aligned}
 out :=_{\delta_2} \text{reduceSeq}(N, \delta_1, \delta_2, F, I, \boxed{E}) &= C(I)_{\delta_2}(\lambda \text{init}. \\
 &\quad \boxed{C(E)}_{N, \delta_1}(\lambda x. \\
 &\quad out :=_{\delta_1} \text{init}; \\
 &\quad \text{for } N(\lambda i. out :=_{\delta_2} F(\boxed{x[i]}, out));))
 \end{aligned}$$

- Start translation from `rise.out`
- Traverse program back to front and lower patterns as specified by their arguments

Lowering to rise.codegen: dot

```

1 func @dot(%out:memref<1xf32>, %a:memref<1024xf32>,
           %b:memref<1024xf32>) {
2   %lhs = rise.in %a : !rise.array<1024, scalar<f32>>
3   %rhs = rise.in %b : !rise.array<1024, scalar<f32>>
4
5   %zip = rise.zip #rise.nat<1024> #rise.scalar<f32>
           #rise.scalar<f32>
6   %zipped = rise.apply %zip, %lhs, %rhs
7   %f = rise.lambda (%tuple, %acc) → !rise.scalar<f32> {
8     %fstFun = rise.fst #rise.scalar<f32> #rise.scalar<f32>
9     %sndFun = rise.snd #rise.scalar<f32> #rise.scalar<f32>
10    %fst = rise.apply %fstFun, %tuple
11    %snd = rise.apply %sndFun, %tuple
12    %result = rise.embed(%fst, %snd, %acc) {
13      %product = mulf %fst, %snd : f32
14      %res = addf %product, %acc : f32
15      return res : f32
16    } : !rise.scalar<f32>
17    rise.return %result : !rise.scalar<f32>
18  }
19  %init = rise.literal #rise.lit<0.0>
20  %reduce = rise.reduceSeq #rise.nat<1024>
           #rise.tuple<scalar<f32>, scalar<f32>> #rise.scalar<f32>
21  %result = rise.apply %reduce, %f, %init, %zipped
22  rise.out %out ← %result
23  return
24 }

```

RISE

lowering
→
lambda{ ... }

```

1 func @dot(%out:memref<1xf32>,%a:memref<1024xf32>,%b:memref<1024xf32>) {
2   %cst0 = constant 0.000000e+00 : f32
3   %c0 = constant 0 : index
4   %c1024 = constant 1024 : index
5   %c1 = constant 1 : index
6   %0 = rise.codegen.zip(%a, %b)
           : (memref<1024xf32>, memref<1024xf32>) → memref<1024xf32>
7   %1 = rise.codegen.idx(%out, %c0) : (memref<1xf32>, index) → f32
8   rise.codegen.assign(%cst0, %1) : (f32, f32) → ()
9   loop.for %i = %c0 to %c1024 step %c1 {
10    %2 = rise.codegen.idx(%0, %i) : (memref<1024xf32>, index) → f32
11    %3 = rise.codegen.idx(%out, %c0) : (memref<1xf32>, index) → f32
12    %4 = rise.codegen.fst(%2) : (f32) → f32
13    %5 = rise.codegen.snd(%2) : (f32) → f32
14    %6 = mulf %4, %5 : f32
15    %7 = addf %6, %3 : f32
16    rise.codegen.assign(%7, %3) : (f32, f32) → ()
17  }
18  return
19 }

```

Loop + RISE Intermediate

$$\begin{aligned}
 out :=_{\delta_2} \text{reduceSeq}(N, \delta_1, \delta_2, \boxed{F} I, E) &= C(\llbracket I \rrbracket)_{\delta_2} (\lambda \text{init}. \\
 &C(\llbracket E \rrbracket)_{N, \delta_1} (\lambda x. \\
 &\quad out :=_{\delta_1} \text{init}; \\
 &\quad \text{for } N(\lambda i. out :=_{\delta_2} F(x[i], out));))
 \end{aligned}$$

- Start translation from rise.out
- Traverse program back to front and lower patterns as specified by their arguments

Lowering to rise.codegen: dot

```

1 func @dot(%out:memref<1xf32>, %a:memref<1024xf32>,
           %b:memref<1024xf32>) {
2   %lhs = rise.in %a : !rise.array<1024, scalar<f32>>
3   %rhs = rise.in %b : !rise.array<1024, scalar<f32>>
4
5   %zip = rise.zip #rise.nat<1024> #rise.scalar<f32>
           #rise.scalar<f32>
6   %zipped = rise.apply %zip, %lhs, %rhs
7   %f = rise.lambda (%tuple, %acc) → !rise.scalar<f32> {
8     %fstFun = rise.fst #rise.scalar<f32> #rise.scalar<f32>
9     %sndFun = rise.snd #rise.scalar<f32> #rise.scalar<f32>
10    %fst = rise.apply %fstFun, %tuple
11    %snd = rise.apply %sndFun, %tuple
12    %result = rise.embed(%fst, %snd, %acc) {
13      %product = mulf %fst, %snd : f32
14      %res = addf %product, %acc : f32
15      return res : f32
16    } : !rise.scalar<f32>
17    rise.return %result : !rise.scalar<f32>
18  }
19  %init = rise.literal #rise.lit<0.0>
20  %reduce = rise.reduceSeq #rise.nat<1024>
           #rise.tuple<scalar<f32>, scalar<f32>> #rise.scalar<f32>
21  %result = rise.apply %reduce, %f, %init, %zipped
22  rise.out %out ← %result
23  return
24 }

```

RISE

lowering
→
lambda{ ... }

```

1 func @dot(%out:memref<1xf32>,%a:memref<1024xf32>,%b:memref<1024xf32>) {
2   %cst0 = constant 0.000000e+00 : f32
3   %c0 = constant 0 : index
4   %c1024 = constant 1024 : index
5   %c1 = constant 1 : index
6   %0 = rise.codegen.zip(%a, %b)
           : (memref<1024xf32>, memref<1024xf32>) → memref<1024xf32>
7   %1 = rise.codegen.idx(%out, %c0) : (memref<1xf32>, index) → f32
8   rise.codegen.assign(%cst0, %1) : (f32, f32) → ()
9   loop.for %i = %c0 to %c1024 step %c1 {
10    %2 = rise.codegen.idx(%0, %i) : (memref<1024xf32>, index) → f32
11    %3 = rise.codegen.idx(%out, %c0) : (memref<1xf32>, index) → f32
12    %4 = rise.codegen.fst(%2) : (f32) → f32
13    %5 = rise.codegen.snd(%2) : (f32) → f32
14    %6 = mulf %4, %5 : f32
15    %7 = addf %6, %3 : f32
16    rise.codegen.assign(%7, %3) : (f32, f32) → ()
17  }
18  return
19 }

```

Loop + RISE Intermediate

$$\begin{aligned}
 out :=_{\delta_2} \text{reduceSeq}(N, \delta_1, \delta_2, \boxed{F} I, E) &= C(\boxed{I})_{\delta_2}(\lambda \text{init}. \\
 &C(\boxed{E})_{N, \delta_1}(\lambda x. \\
 &\quad out :=_{\delta_1} \text{init}; \\
 &\quad \text{for } N(\lambda i. out :=_{\delta_2} F(x[i], out));))
 \end{aligned}$$

- Start translation from rise.out
- Traverse program back to front and lower patterns as specified by their arguments

Lowering to rise.codegen: dot

```

1 func @dot(%out:memref<1xf32>, %a:memref<1024xf32>,
           %b:memref<1024xf32>) {
2   %lhs = rise.in %a : !rise.array<1024, scalar<f32>>
3   %rhs = rise.in %b : !rise.array<1024, scalar<f32>>
4
5   %zip = rise.zip #rise.nat<1024> #rise.scalar<f32>
           #rise.scalar<f32>
6   %zipped = rise.apply %zip, %lhs, %rhs
7   %f = rise.lambda (%tuple, %acc) → !rise.scalar<f32> {
8     %fstFun = rise.fst #rise.scalar<f32> #rise.scalar<f32>
9     %sndFun = rise.snd #rise.scalar<f32> #rise.scalar<f32>
10    %fst = rise.apply %fstFun, %tuple
11    %snd = rise.apply %sndFun, %tuple
12    %result = rise.embed(%fst, %snd, %acc) {
13      %product = mulf %fst, %snd : f32
14      %res = addf %product, %acc : f32
15      return res : f32
16    } : !rise.scalar<f32>
17    rise.return %result : !rise.scalar<f32>
18  }
19  %init = rise.literal #rise.lit<0.0>
20  %reduce = rise.reduceSeq #rise.nat<1024>
           #rise.tuple<scalar<f32>, scalar<f32>> #rise.scalar<f32>
21  %result = rise.apply %reduce, %f, %init, %zipped
22  rise.out %out ← %result
23  return
24 }

```

RISE

lowering
→
lambda{ ... }

```

1 func @dot(%out:memref<1xf32>,%a:memref<1024xf32>,%b:memref<1024xf32>) {
2   %cst0 = constant 0.000000e+00 : f32
3   %c0 = constant 0 : index
4   %c1024 = constant 1024 : index
5   %c1 = constant 1 : index
6   %0 = rise.codegen.zip(%a, %b)
           : (memref<1024xf32>, memref<1024xf32>) → memref<1024xf32>
7   %1 = rise.codegen.idx(%out, %c0) : (memref<1xf32>, index) → f32
8   rise.codegen.assign(%cst0, %1) : (f32, f32) → ()
9   loop.for %i = %c0 to %c1024 step %c1 {
10    %2 = rise.codegen.idx(%0, %i) : (memref<1024xf32>, index) → f32
11    %3 = rise.codegen.idx(%out, %c0) : (memref<1xf32>, index) → f32
12    %4 = rise.codegen.fst(%2) : (f32) → f32
13    %5 = rise.codegen.snd(%2) : (f32) → f32
14    %6 = mulf %4, %5 : f32
15    %7 = addf %6, %3 : f32
16    rise.codegen.assign(%7, %3) : (f32, f32) → ()
17  }
18  return
19 }

```

Loop + RISE Intermediate

$$\begin{aligned}
 out :=_{\delta_2} \text{reduceSeq}(N, \delta_1, \delta_2, \boxed{F}, I, E) &= C(\llbracket I \rrbracket)_{\delta_2}(\lambda \text{init}. \\
 &C(\llbracket E \rrbracket)_{N, \delta_1}(\lambda x. \\
 &\quad out :=_{\delta_1} \text{init}; \\
 &\quad \text{for } N(\lambda i. out :=_{\delta_2} F(x[i], out));))
 \end{aligned}$$

- Start translation from `rise.out`
- Traverse program back to front and lower patterns as specified by their arguments

Lowering to rise.codegen: dot

```

1 func @dot(%out:memref<1xf32>, %a:memref<1024xf32>,
           %b:memref<1024xf32>) {
2   %lhs = rise.in %a : !rise.array<1024, scalar<f32>>
3   %rhs = rise.in %b : !rise.array<1024, scalar<f32>>
4
5   %zip = rise.zip #rise.nat<1024> #rise.scalar<f32>
           #rise.scalar<f32>
6   %zipped = rise.apply %zip, %lhs, %rhs
7   %f = rise.lambda (%tuple, %acc) → !rise.scalar<f32> {
8     %fstFun = rise.fst #rise.scalar<f32> #rise.scalar<f32>
9     %sndFun = rise.snd #rise.scalar<f32> #rise.scalar<f32>
10    %fst = rise.apply %fstFun, %tuple
11    %snd = rise.apply %sndFun, %tuple
12    %result = rise.embed(%fst, %snd, %acc) {
13      %product = mulf %fst, %snd : f32
14      %res = addf %product, %acc : f32
15      return res : f32
16    } : !rise.scalar<f32>
17    rise.return %result : !rise.scalar<f32>
18  }
19 %init = rise.literal #rise.lit<0.0>
20 %reduce = rise.reduceSeq #rise.nat<1024>
           #rise.tuple<scalar<f32>, scalar<f32>> #rise.scalar<f32>
21 %result = rise.apply %reduce, %f, %init, %zipped
22 rise.out %out ← %result
23 return
24 }

```

RISE

lowering
embed{ ... }

```

1 func @dot(%out:memref<1xf32>,%a:memref<1024xf32>,%b:memref<1024xf32>) {
2   %cst0 = constant 0.000000e+00 : f32
3   %c0 = constant 0 : index
4   %c1024 = constant 1024 : index
5   %c1 = constant 1 : index
6   %0 = rise.codegen.zip(%a, %b)
           : (memref<1024xf32>, memref<1024xf32>) → memref<1024xf32>
7   %1 = rise.codegen.idx(%out, %c0) : (memref<1xf32>, index) → f32
8   rise.codegen.assign(%cst0, %1) : (f32, f32) → ()
9   loop.for %i = %c0 to %c1024 step %c1 {
10    %2 = rise.codegen.idx(%0, %i) : (memref<1024xf32>, index) → f32
11    %3 = rise.codegen.idx(%out, %c0) : (memref<1xf32>, index) → f32
12    %4 = rise.codegen.fst(%2) : (f32) → f32
13    %5 = rise.codegen.snd(%2) : (f32) → f32
14    { %6 = mulf %4, %5 : f32
15      %7 = addf %6, %3 : f32
16      rise.codegen.assign(%7, %3) : (f32, f32) → ()
17    }
18    return
19 }

```

Loop + RISE Intermediate

$$\begin{aligned}
 out :=_{\delta_2} \text{reduceSeq}(N, \delta_1, \delta_2, \boxed{F} I, E) &= C(\llbracket I \rrbracket)_{\delta_2} (\lambda \text{init}. \\
 &C(\llbracket E \rrbracket)_{N, \delta_1} (\lambda x. \\
 &\quad out :=_{\delta_1} \text{init}; \\
 &\quad \text{for } N(\lambda i. out :=_{\delta} F(x[i], out));))
 \end{aligned}$$

- Start translation from `rise.out`
- Traverse program back to front and lower patterns as specified by their arguments

Lowering to rise.codegen: dot

```

1 func @dot(%out:memref<1xf32>, %a:memref<1024xf32>,
           %b:memref<1024xf32>) {
2   %lhs = rise.in %a : !rise.array<1024, scalar<f32>>
3   %rhs = rise.in %b : !rise.array<1024, scalar<f32>>
4
5   %zip = rise.zip #rise.nat<1024> #rise.scalar<f32>
           #rise.scalar<f32>
6   %zipped = rise.apply %zip, %lhs, %rhs
7   %f = rise.lambda (%tuple, %acc) → !rise.scalar<f32> {
8     %fstFun = rise.fst #rise.scalar<f32> #rise.scalar<f32>
9     %sndFun = rise.snd #rise.scalar<f32> #rise.scalar<f32>
10    %fst = rise.apply %fstFun, %tuple
11    %snd = rise.apply %sndFun, %tuple
12    %result = rise.embed(%fst, %snd, %acc) {
13      %product = mulf %fst, %snd : f32
14      %res = addf %product, %acc : f32
15      return res : f32
16    } : !rise.scalar<f32>
17    rise.return %result : !rise.scalar<f32>
18  }
19  %init = rise.literal #rise.lit<0.0>
20  %reduce = rise.reduceSeq #rise.nat<1024>
           #rise.tuple<scalar<f32>, scalar<f32>> #rise.scalar<f32>
21  %result = rise.apply %reduce, %f, %init, %zipped
22  rise.out %out ← %result
23  return
24 }

```

RISE

lowering

```

1 func @dot(%out:memref<1xf32>,%a:memref<1024xf32>,%b:memref<1024xf32>) {
2   %cst0 = constant 0.000000e+00 : f32
3   %c0 = constant 0 : index
4   %c1024 = constant 1024 : index
5   %c1 = constant 1 : index
6   %0 = rise.codegen.zip(%a, %b)
           : (memref<1024xf32>, memref<1024xf32>) → memref<1024xf32>
7   %1 = rise.codegen.idx(%out, %c0) : (memref<1xf32>, index) → f32
8   rise.codegen.assign(%cst0, %1) : (f32, f32) → ()
9   loop.for %i = %c0 to %c1024 step %c1 {
10    %2 = rise.codegen.idx(%0, %i) : (memref<1024xf32>, index) → f32
11    %3 = rise.codegen.idx(%out, %c0) : (memref<1xf32>, index) → f32
12    %4 = rise.codegen.fst(%2) : (f32) → f32
13    %5 = rise.codegen.snd(%2) : (f32) → f32
14    %6 = mulf %4, %5 : f32
15    %7 = addf %6, %3 : f32
16    rise.codegen.assign(%7, %3) : (f32, f32) → ()
17  }
18  return
19 }

```

Loop + RISE Intermediate

$$\begin{aligned}
 out :=_{\delta_2} \text{reduceSeq}(N, \delta_1, \delta_2, F, I, E) &= C(I)_{\delta_2}(\lambda \text{init}. \\
 &C(E)_{N, \delta_1}(\lambda x. \\
 &\quad out :=_{\delta_1} \text{init}; \\
 &\quad \text{for } N (\lambda i. out :=_{\delta_2} F(x[i], out));))
 \end{aligned}$$

- Start translation from `rise.out`
- Traverse program back to front and lower patterns as specified by their arguments

Lowering of the RISE dialect

Lowering of patterns is not context free, e.g. zip influences the code generation of the next pattern

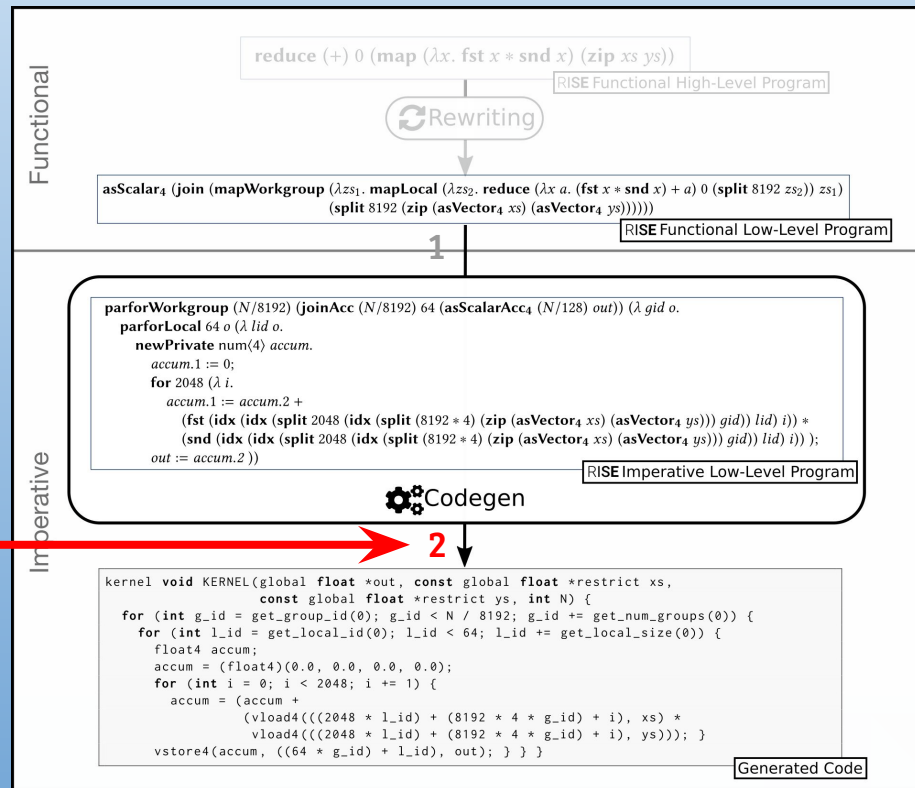
Lowering is divided into two steps:

1. Lowering functional to intermediate imperative

2. Lowering to final target code

Lowering described formally in:

Strategy Preserving Compilation for Parallel Functional Code
by Robert Atkey, Michel Steuwer, Sam Lindley, and Christophe Dubach
<https://arxiv.org/abs/1710.08332>



Lowering rise.codegen to loop: Matrix Multiplication

```
func @mm(%out, %A, %B) {  
  // map(map(dot))  
}
```

lowering

```
1 func @mm(%out, %A, %B) {  
2   %cst0 = constant 0.000000e+00 : f32  
3   %c0 = constant 0 : index  
4   %c1024 = constant 1024 : index  
5   %c1 = constant 1 : index  
6   loop.for %i = %c0 to %c1024 step %c1 {  
7     %0 = rise.codegen.idx(%A, %i)  
8     %1 = rise.codegen.idx(%out, %i)  
9     loop.for %j = %c0 to %c1024 step %c1 {  
10      %2 = rise.codegen.idx(%B, %j)  
11      %3 = rise.codegen.idx(%1, %j)  
12      %4 = rise.codegen.zip(%0, %2)  
13      %5 = rise.codegen.idx(%3, %c0)  
14      rise.codegen.assign(%cst0, %5)  
15      loop.for %k = %c0 to %c1024 step %c1 {  
16        %6 = rise.codegen.idx(%4, %k)  
17        %7 = rise.codegen.idx(%3, %c0)  
18        %8 = rise.codegen.fst(%6)  
19        %9 = rise.codegen.snd(%6)  
20        %10 = mulf %8, %9 : f32  
21        %11 = addf %10, %7 : f32  
22        rise.codegen.assign(%11, %7)  
23      }  
24    }  
25  }  
26  return
```

Loop + RISE Intermediate

```
1 func @dot(%out:memref<1xf32>,%a:memref<1024xf32>,%b:memref<1024xf32>) {  
2   %cst0 = constant 0.000000e+00 : f32  
3   %c0 = constant 0 : index  
4   %c1024 = constant 1024 : index  
5   %c1 = constant 1 : index  
6   %0 = rise.codegen.zip(%a, %b)  
7   %1 = rise.codegen.idx(%out, %c0) : (memref<1xf32>, index) → f32  
8   rise.codegen.assign(%cst0, %1) : (f32, f32) → ()  
9   loop.for %i = %c0 to %c1024 step %c1 {  
10    %2 = rise.codegen.idx(%0, %i) : (memref<1024xf32>, index) → f32  
11    %3 = rise.codegen.idx(%out, %c0) : (memref<1xf32>, index) → f32  
12    %4 = rise.codegen.fst(%2) : (f32) → f32  
13    %5 = rise.codegen.snd(%2) : (f32) → f32  
14    %6 = mulf %4, %5 : f32  
15    %7 = addf %6, %3 : f32  
16    rise.codegen.assign(%7, %3) : (f32, f32) → ()  
17  }  
18  return  
19 }
```

Loop + RISE Intermediate

- Matrix multiply builds on dot -> similar inner structure
- Only differences:
 - different **input** to **rise.codegen.zip**
 - accumulating **indexed** with j and i

Lowering rise.codegen to loop: Matrix Multiplication

```
1 func @mm(%out: memref<1024×1024xf32>, %A: memref<1024×1024xf32>, %B: memref<1024×1024xf32>) {  
    %cst0 = constant 0.000000e+00 : f32  
    %c0 = constant 0 : index  
    %c1024 = constant 1024 : index  
    %c1 = constant 1 : index  
    loop.for %i = %c0 to %c1024 step %c1 {  
        %0 = rise.codegen.idx(%A, %i) : (memref<1024×1024xf32>, index) → memref<1024xf32>  
        %1 = rise.codegen.idx(%out, %i) : (memref<1024×1024xf32>, index) → memref<1024xf32>  
        loop.for %j = %c0 to %c1024 step %c1 {  
            %2 = rise.codegen.idx(%B, %j) : (memref<1024×1024xf32>, index) → memref<1024xf32>  
            %3 = rise.codegen.idx(%1, %j) : (memref<1024xf32>, index) → memref<1024xf32>  
            %4 = rise.codegen.zip(%0, %2) : (memref<1024xf32>, memref<1024xf32>) → memref<1024xf32>  
            %5 = rise.codegen.idx(%3, %c0) : (memref<1024xf32>, index) → f32  
            rise.codegen.assign(%cst0, %5) : (f32, f32) → ()  
            loop.for %k = %c0 to %c1024 step %c1 {  
                %6 = rise.codegen.idx(%4, %k) : (memref<1024xf32>, index) → f32  
                %7 = rise.codegen.idx(%3, %c0) : (memref<1024xf32>, index) → f32  
                %8 = rise.codegen.fst(%6) : (f32) → f32  
                %9 = rise.codegen.snd(%6) : (f32) → f32  
                %10 = mulf %8, %9 : f32  
                %11 = addf %10, %7 : f32  
                rise.codegen.assign(%11, %7) : (f32, f32) → ()  
            }  
        }  
    }  
    return  
}
```

Loop + RISE Intermediate

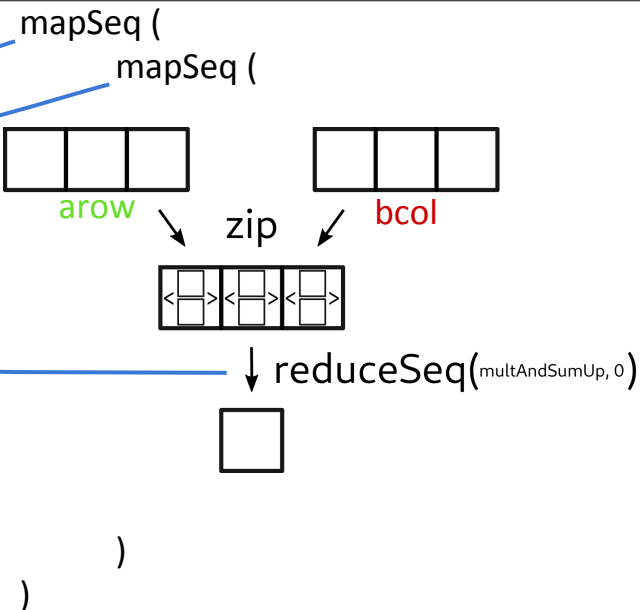
Lowering rise.codegen to loop: Matrix Multiplication

```

1 func @mm(%out: memref<1024×1024xf32>, %A: memref<1024×1024xf32>, %B: memref<1024×1024xf32>) {
2   %cst0 = constant 0.000000e+00 : f32
3   %c0 = constant 0 : index
4   %c1024 = constant 1024 : index
5   %c1 = constant 1 : index
6   loop.for %i = %c0 to %c1024 step %c1 {
7     %0 = rise.codegen.idx(%A, %i) : (memref<1024×1024xf32>, index) → memref<1024xf32>
8     %1 = rise.codegen.idx(%out, %i) : (memref<1024×1024xf32>, index) → memref<1024xf32>
9     loop.for %j = %c0 to %c1024 step %c1 {
10      %2 = rise.codegen.idx(%B, %j) : (memref<1024×1024xf32>, index) → memref<1024xf32>
11      %3 = rise.codegen.idx(%1, %j) : (memref<1024xf32>, index) → memref<1024xf32>
12      %4 = rise.codegen.zip(%0, %2) : (memref<1024xf32>, memref<1024xf32>) → memref<1024xf32>
13      %5 = rise.codegen.idx(%3, %c0) : (memref<1024xf32>, index) → f32
14      rise.codegen.assign(%cst0, %5) : (f32, f32) → ()
15      loop.for %k = %c0 to %c1024 step %c1 {
16        %6 = rise.codegen.idx(%4, %k) : (memref<1024xf32>, index) → f32
17        %7 = rise.codegen.idx(%3, %c0) : (memref<1024xf32>, index) → f32
18        %8 = rise.codegen.fst(%6) : (f32) → f32
19        %9 = rise.codegen.snd(%6) : (f32) → f32
20        %10 = mulf %8, %9 : f32
21        %11 = addf %10, %7 : f32
22        rise.codegen.assign(%11, %7) : (f32, f32) → ()
23      }
24    }
25  }
26 }

```

Loop + RISE Intermediate



Lowering rise.codegen to loop: Matrix Multiplication

```
1 func @mm(%out: memref<1024×1024xf32>, %A: memref<1024×1024xf32>, %B: memref<1024×1024xf32>) {
2   %cst0 = constant 0.000000e+00 : f32
3   %c0 = constant 0 : index
4   %c1024 = constant 1024 : index
5   %c1 = constant 1 : index
6   loop.for %i = %c0 to %c1024 step %c1 {
7     %0 = rise.codegen.idx(%A, %i) : (memref<1024×1024xf32>, index) → memref<1024xf32>
8     %1 = rise.codegen.idx(%out, %i) : (memref<1024×1024xf32>, index) → memref<1024xf32>
9     loop.for %j = %c0 to %c1024 step %c1 {
10      %2 = rise.codegen.idx(%B, %j) : (memref<1024×1024xf32>, index) → memref<1024xf32>
11      %3 = rise.codegen.idx(%1, %j) : (memref<1024xf32>, index) → memref<1024xf32>
12      %4 = rise.codegen.zip(%0, %2) : (memref<1024xf32>, memref<1024xf32>) → memref<1024xf32>
13      %5 = rise.codegen.idx(%3, %c0) : (memref<1024xf32>, index) → f32
14      rise.codegen.assign(%cst0, %5) : (f32, f32) → ()
15      loop.for %k = %c0 to %c1024 step %c1 {
16        %6 = rise.codegen.idx(%4, %k) : (memref<1024xf32>, index) → f32
17        %7 = rise.codegen.idx(%3, %c0) : (memref<1024xf32>, index) → f32
18        %8 = rise.codegen.fst(%6) : (f32) → f32
19        %9 = rise.codegen.snd(%6) : (f32) → f32
20        %10 = mulf %8, %9 : f32
21        %11 = addf %10, %7 : f32
22        rise.codegen.assign(%11, %7) : (f32, f32) → ()
23      }
24    }
25  }
26  return
27 }
```

Loop + RISE Intermediate

? = ?

- Start codegen from rise.codegen.assign

Lowering rise.codegen to loop: Matrix Multiplication



```
1 func @mm(%out: memref<1024x1024xf32>, %A: memref<1024x1024xf32>, %B: memref<1024x1024xf32>) {
2   %cst0 = constant 0.000000e+00 : f32
3   %c0 = constant 0 : index
4   %c1024 = constant 1024 : index
5   %c1 = constant 1 : index
6   loop.for %i = %c0 to %c1024 step %c1 {
7     %0 = rise.codegen.idx(%A, %i) : (memref<1024x1024xf32>, index) → memref<1024xf32>
8     %1 = rise.codegen.idx(%out, %i) : (memref<1024x1024xf32>, index) → memref<1024xf32>
9     loop.for %j = %c0 to %c1024 step %c1 {
10      %2 = rise.codegen.idx(%B, %j) : (memref<1024x1024xf32>, index) → memref<1024xf32>
11      %3 = rise.codegen.idx(%1, %j) : (memref<1024xf32>, index) → memref<1024xf32>
12      %4 = rise.codegen.zip(%0, %2) : (memref<1024xf32>, memref<1024xf32>) → memref<1024xf32>
13      %5 = rise.codegen.idx(%3, %c0) : (memref<1024xf32>, index) → f32
14      rise.codegen.assign(%cst0, %5) : (f32, f32) → ()
15      loop.for %k = %c0 to %c1024 step %c1 {
16        %6 = rise.codegen.idx(%4, %k) : (memref<1024xf32>, index) → f32
17        %7 = rise.codegen.idx(%3, %c0) : (memref<1024xf32>, index) → f32
18        %8 = rise.codegen.fst(%6) : (f32) → f32
19        %9 = rise.codegen.snd(%6) : (f32) → f32
20        %10 = mulf %8, %9 : f32
21        → %11 = addf %10, %7 : f32
22        rise.codegen.assign(%11, %7) : (f32, f32) → ()
23      }
24    }
25  }
26  return
27 }
```

Loop + RISE Intermediate

? = ? + ?

- Start codegen from `rise.codegen.assign`
- Traverse program back to front following references

Lowering rise.codegen to loop: Matrix Multiplication

```
1 func @mm(%out: memref<1024x1024xf32>, %A: memref<1024x1024xf32>, %B: memref<1024x1024xf32>) {
2   %cst0 = constant 0.000000e+00 : f32
3   %c0 = constant 0 : index
4   %c1024 = constant 1024 : index
5   %c1 = constant 1 : index
6   loop.for %i = %c0 to %c1024 step %c1 {
7     %0 = rise.codegen.idx(%A, %i) : (memref<1024x1024xf32>, index) → memref<1024xf32>
8     %1 = rise.codegen.idx(%out, %i) : (memref<1024x1024xf32>, index) → memref<1024xf32>
9     loop.for %j = %c0 to %c1024 step %c1 {
10      %2 = rise.codegen.idx(%B, %j) : (memref<1024x1024xf32>, index) → memref<1024xf32>
11      %3 = rise.codegen.idx(%1, %j) : (memref<1024xf32>, index) → memref<1024xf32>
12      %4 = rise.codegen.zip(%0, %2) : (memref<1024xf32>, memref<1024xf32>) → memref<1024xf32>
13      %5 = rise.codegen.idx(%3, %c0) : (memref<1024xf32>, index) → f32
14      rise.codegen.assign(%cst0, %5) : (f32, f32) → ()
15      loop.for %k = %c0 to %c1024 step %c1 {
16        %6 = rise.codegen.idx(%4, %k) : (memref<1024xf32>, index) → f32
17        %7 = rise.codegen.idx(%3, %c0) : (memref<1024xf32>, index) → f32
18        %8 = rise.codegen.fst(%6) : (f32) → f32
19        %9 = rise.codegen.snd(%6) : (f32) → f32
20         %10 = mulf %8, %9 : f32
21         %11 = addf %10, %7 : f32
22        rise.codegen.assign(%11, %7) : (f32, f32) → ()
23      }
24    }
25  }
26  return
27 }
```

Loop + RISE Intermediate

? = ? * ? + ?

- Start codegen from `rise.codegen.assign`
- Traverse program back to front following references

Lowering rise.codegen to loop: Matrix Multiplication

```
1 func @mm(%out: memref<1024x1024xf32>, %A: memref<1024x1024xf32>, %B: memref<1024x1024xf32>) {
2   %cst0 = constant 0.000000e+00 : f32
3   %c0 = constant 0 : index
4   %c1024 = constant 1024 : index
5   %c1 = constant 1 : index
6   loop.for %i = %c0 to %c1024 step %c1 {
7     %0 = rise.codegen.idx(%A, %i) : (memref<1024x1024xf32>, index) → memref<1024xf32>
8     %1 = rise.codegen.idx(%out, %i) : (memref<1024x1024xf32>, index) → memref<1024xf32>
9     loop.for %j = %c0 to %c1024 step %c1 {
10      %2 = rise.codegen.idx(%B, %j) : (memref<1024x1024xf32>, index) → memref<1024xf32>
11      %3 = rise.codegen.idx(%1, %j) : (memref<1024xf32>, index) → memref<1024xf32>
12      %4 = rise.codegen.zip(%0, %2) : (memref<1024xf32>, memref<1024xf32>) → memref<1024xf32>
13      %5 = rise.codegen.idx(%3, %c0) : (memref<1024xf32>, index) → f32
14      rise.codegen.assign(%cst0, %5) : (f32, f32) → ()
15      loop.for %k = %c0 to %c1024 step %c1 {
16        %6 = rise.codegen.idx(%4, %k) : (memref<1024xf32>, index) → f32
17        %7 = rise.codegen.idx(%3, %c0) : (memref<1024xf32>, index) → f32
18        %8 = rise.codegen.fst(%6) : (f32) → f32
19        %9 = rise.codegen.snd(%6) : (f32) → f32
20        %10 = mulf %8, %9 : f32
21        %11 = addf %10, %7 : f32
22        rise.codegen.assign(%11, %7) : (f32, f32) → ()
23      }
24    }
25  }
26  return
27 }
```

Path:

fst

Loop + RISE Intermediate

? = ? * ? + ?

- Start codegen from rise.codegen.assign
- Traverse program back to front following references
- build up path

Lowering rise.codegen to loop: Matrix Multiplication

```
1 func @mm(%out: memref<1024×1024xf32>, %A: memref<1024×1024xf32>, %B: memref<1024×1024xf32>) {
2   %cst0 = constant 0.000000e+00 : f32
3   %c0 = constant 0 : index
4   %c1024 = constant 1024 : index
5   %c1 = constant 1 : index
6   loop.for %i = %c0 to %c1024 step %c1 {
7     %0 = rise.codegen.idx(%A, %i) : (memref<1024×1024xf32>, index) → memref<1024xf32>
8     %1 = rise.codegen.idx(%out, %i) : (memref<1024×1024xf32>, index) → memref<1024xf32>
9     loop.for %j = %c0 to %c1024 step %c1 {
10      %2 = rise.codegen.idx(%B, %j) : (memref<1024×1024xf32>, index) → memref<1024xf32>
11      %3 = rise.codegen.idx(%1, %j) : (memref<1024xf32>, index) → memref<1024xf32>
12      %4 = rise.codegen.zip(%0, %2) : (memref<1024xf32>, memref<1024xf32>) → memref<1024xf32>
13      %5 = rise.codegen.idx(%3, %c0) : (memref<1024xf32>, index) → f32
14      rise.codegen.assign(%cst0, %5) : (f32, f32) → ()
15      loop.for %k = %c0 to %c1024 step %c1 {
16        %6 = rise.codegen.idx(%4, %k) : (memref<1024xf32>, index) → f32
17        %7 = rise.codegen.idx(%3, %c0) : (memref<1024xf32>, index) → f32
18        %8 = rise.codegen.fst(%6) : (f32) → f32
19        %9 = rise.codegen.snd(%6) : (f32) → f32
20        %10 = mulf %8, %9 : f32
21        %11 = addf %10, %7 : f32
22        rise.codegen.assign(%11, %7) : (f32, f32) → ()
23      }
24    }
25  }
26  return
27 }
```

Path:

idx(%k)
fst

Loop + RISE Intermediate

? = ?[?] * ? + ?

- Start codegen from rise.codegen.assign
- Traverse program back to front following references
- build up path

Lowering rise.codegen to loop: Matrix Multiplication

```
1 func @mm(%out: memref<1024×1024xf32>, %A: memref<1024×1024xf32>, %B: memref<1024×1024xf32>) {
2   %cst0 = constant 0.000000e+00 : f32
3   %c0 = constant 0 : index
4   %c1024 = constant 1024 : index
5   %c1 = constant 1 : index
6   loop.for %i = %c0 to %c1024 step %c1 {
7     %0 = rise.codegen.idx(%A, %i) : (memref<1024×1024xf32>, index) → memref<1024xf32>
8     %1 = rise.codegen.idx(%out, %i) : (memref<1024×1024xf32>, index) → memref<1024xf32>
9     loop.for %j = %c0 to %c1024 step %c1 {
10      %2 = rise.codegen.idx(%B, %j) : (memref<1024×1024xf32>, index) → memref<1024xf32>
11      %3 = rise.codegen.idx(%1, %j) : (memref<1024xf32>, index) → memref<1024xf32>
12      %4 = rise.codegen.zip(%0, %2) : (memref<1024xf32>, memref<1024xf32>) → memref<1024xf32>
13      %5 = rise.codegen.idx(%3, %c0) : (memref<1024xf32>, index) → f32
14      rise.codegen.assign(%cst0, %5) : (f32, f32) → ()
15      loop.for %k = %c0 to %c1024 step %c1 {
16        %6 = rise.codegen.idx(%4, %k) : (memref<1024xf32>, index) → f32
17        %7 = rise.codegen.idx(%3, %c0) : (memref<1024xf32>, index) → f32
18        %8 = rise.codegen.fst(%6) : (f32) → f32
19        %9 = rise.codegen.snd(%6) : (f32) → f32
20        %10 = mulf %8, %9 : f32
21        %11 = addf %10, %7 : f32
22        rise.codegen.assign(%11, %7) : (f32, f32) → ()
23      }
24    }
25  }
26  return
27 }
```

Path:

idx(%k)
~~fst~~

codegen(lhs)

Loop + RISE Intermediate

? = ?[?] * ? + ?

- Start codegen from rise.codegen.assign
- Traverse program back to front following references
- build up path

Lowering rise.codegen to loop: Matrix Multiplication

```
1 func @mm(%out: memref<1024×1024xf32>, %A: memref<1024×1024xf32>, %B: memref<1024×1024xf32>) {
2   %cst0 = constant 0.000000e+00 : f32
3   %c0 = constant 0 : index
4   %c1024 = constant 1024 : index
5   %c1 = constant 1 : index
6   loop.for %i = %c0 to %c1024 step %c1 {
7     %0 = rise.codegen.idx(%A, %i) : (memref<1024×1024xf32>, index) → memref<1024xf32>
8     %1 = rise.codegen.idx(%out, %i) : (memref<1024×1024xf32>, index) → memref<1024xf32>
9     loop.for %j = %c0 to %c1024 step %c1 {
10      %2 = rise.codegen.idx(%B, %j) : (memref<1024×1024xf32>, index) → memref<1024xf32>
11      %3 = rise.codegen.idx(%1, %j) : (memref<1024xf32>, index) → memref<1024xf32>
12      %4 = rise.codegen.zip(%0, %2) : (memref<1024xf32>, memref<1024xf32>) → memref<1024xf32>
13      %5 = rise.codegen.idx(%3, %c0) : (memref<1024xf32>, index) → f32
14      rise.codegen.assign(%cst0, %5) : (f32, f32) → ()
15      loop.for %k = %c0 to %c1024 step %c1 {
16        %6 = rise.codegen.idx(%4, %k) : (memref<1024xf32>, index) → f32
17        %7 = rise.codegen.idx(%3, %c0) : (memref<1024xf32>, index) → f32
18        %8 = rise.codegen.fst(%6) : (f32) → f32
19        %9 = rise.codegen.snd(%6) : (f32) → f32
20        %10 = mulf %8, %9 : f32
21        %11 = addf %10, %7 : f32
22        rise.codegen.assign(%11, %7) : (f32, f32) → ()
23      }
24    }
25  }
26  return
27 }
```

Path:

idx(%i)
idx(%k)

? = ?[?,?] * ? + ?

Loop + RISE Intermediate

- Start codegen from rise.codegen.assign
- Traverse program back to front following references
- build up path

Lowering rise.codegen to loop: Matrix Multiplication

```
1 func @mm(%out: memref<1024×1024xf32>, %A: memref<1024×1024xf32>, %B: memref<1024×1024xf32>) {  
2   %cst0 = constant 0.000000e+00 : f32  
3   %c0 = constant 0 : index  
4   %c1024 = constant 1024 : index  
5   %c1 = constant 1 : index  
6   loop.for %i = %c0 to %c1024 step %c1 {  
7     %0 = rise.codegen.idx(%A, %i) : (memref<1024×1024xf32>, index) → memref<1024xf32>  
8     %1 = rise.codegen.idx(%out, %i) : (memref<1024×1024xf32>, index) → memref<1024xf32>  
9     loop.for %j = %c0 to %c1024 step %c1 {  
10      %2 = rise.codegen.idx(%B, %j) : (memref<1024×1024xf32>, index) → memref<1024xf32>  
11      %3 = rise.codegen.idx(%1, %j) : (memref<1024xf32>, index) → memref<1024xf32>  
12      %4 = rise.codegen.zip(%0, %2) : (memref<1024xf32>, memref<1024xf32>) → memref<1024xf32>  
13      %5 = rise.codegen.idx(%3, %c0) : (memref<1024xf32>, index) → f32  
14      rise.codegen.assign(%cst0, %5) : (f32, f32) → ()  
15      loop.for %k = %c0 to %c1024 step %c1 {  
16        %6 = rise.codegen.idx(%4, %k) : (memref<1024xf32>, index) → f32  
17        %7 = rise.codegen.idx(%3, %c0) : (memref<1024xf32>, index) → f32  
18        %8 = rise.codegen.fst(%6) : (f32) → f32  
19        %9 = rise.codegen.snd(%6) : (f32) → f32  
20        %10 = mulf %8, %9 : f32  
21        %11 = addf %10, %7 : f32  
22        rise.codegen.assign(%11, %7) : (f32, f32) → ()  
23      }  
24    }  
25  }  
26  return  
27 }
```

Path:

idx(%i)
idx(%k)

? = A[?,?] * ? + ?

Loop + RISE Intermediate

- Start codegen from rise.codegen.assign
- Traverse program back to front following references
- build up path

Lowering rise.codegen to loop: Matrix Multiplication

```
1 func @mm(%out: memref<1024×1024xf32>, %A: memref<1024×1024xf32>, %B: memref<1024×1024xf32>) {
2   %cst0 = constant 0.000000e+00 : f32
3   %c0 = constant 0 : index
4   %c1024 = constant 1024 : index
5   %c1 = constant 1 : index
6   loop.for %i = %c0 to %c1024 step %c1 {
7     %0 = rise.codegen.idx(%A, %i) : (memref<1024×1024xf32>, index) → memref<1024xf32>
8     %1 = rise.codegen.idx(%out, %i) : (memref<1024×1024xf32>, index) → memref<1024xf32>
9     loop.for %j = %c0 to %c1024 step %c1 {
10      %2 = rise.codegen.idx(%B, %j) : (memref<1024×1024xf32>, index) → memref<1024xf32>
11      %3 = rise.codegen.idx(%1, %j) : (memref<1024xf32>, index) → memref<1024xf32>
12      %4 = rise.codegen.zip(%0, %2) : (memref<1024xf32>, memref<1024xf32>) → memref<1024xf32>
13      %5 = rise.codegen.idx(%3, %c0) : (memref<1024xf32>, index) → f32
14      rise.codegen.assign(%cst0, %5) : (f32, f32) → ()
15      loop.for %k = %c0 to %c1024 step %c1 {
16        %6 = rise.codegen.idx(%4, %k) : (memref<1024xf32>, index) → f32
17        %7 = rise.codegen.idx(%3, %c0) : (memref<1024xf32>, index) → f32
18        %8 = rise.codegen.fst(%6) : (f32) → f32
19        %9 = rise.codegen.snd(%6) : (f32) → f32
20        %10 = mulf %8, %9 : f32
21        %11 = addf %10, %7 : f32
22        rise.codegen.assign(%11, %7) : (f32, f32) → ()
23      }
24    }
25  }
26  return
27 }
```

Path:

idx(%i)
idx(%k)

? = A[k,i] * ? + ?

Loop + RISE Intermediate

- Start codegen from rise.codegen.assign
- Traverse program back to front following references
- build up path
- reverse path

Lowering rise.codegen to loop: Matrix Multiplication

```
1 func @mm(%out: memref<1024×1024xf32>, %A: memref<1024×1024xf32>, %B: memref<1024×1024xf32>) {
2   %cst0 = constant 0.000000e+00 : f32
3   %c0 = constant 0 : index
4   %c1024 = constant 1024 : index
5   %c1 = constant 1 : index
6   loop.for %i = %c0 to %c1024 step %c1 {
7     %0 = rise.codegen.idx(%A, %i) : (memref<1024×1024xf32>, index) → memref<1024xf32>
8     %1 = rise.codegen.idx(%out, %i) : (memref<1024×1024xf32>, index) → memref<1024xf32>
9     loop.for %j = %c0 to %c1024 step %c1 {
10      %2 = rise.codegen.idx(%B, %j) : (memref<1024×1024xf32>, index) → memref<1024xf32>
11      %3 = rise.codegen.idx(%1, %j) : (memref<1024xf32>, index) → memref<1024xf32>
12      %4 = rise.codegen.zip(%0, %2) : (memref<1024xf32>, memref<1024xf32>) → memref<1024xf32>
13      %5 = rise.codegen.idx(%3, %c0) : (memref<1024xf32>, index) → f32
14      rise.codegen.assign(%cst0, %5) : (f32, f32) → ()
15      loop.for %k = %c0 to %c1024 step %c1 {
16        %6 = rise.codegen.idx(%4, %k) : (memref<1024xf32>, index) → f32
17        %7 = rise.codegen.idx(%3, %c0) : (memref<1024xf32>, index) → f32
18        %8 = rise.codegen.fst(%6) : (f32) → f32
19        %9 = rise.codegen.snd(%6) : (f32) → f32
20        %10 = mulf %8, %9 : f32
21        %11 = addf %10, %7 : f32
22        rise.codegen.assign(%11, %7) : (f32, f32) → ()
23      }
24    }
25  }
26  return
27 }
```

generate()

%new8 = load %A[%k, %i]

Path:

idx(%i)
idx(%k)

? = A[k,i] * ? + ?

Loop + RISE Intermediate

- Start codegen from `rise.codegen.assign`
- Traverse program back to front following references
- build up path
- reverse path and **generate** new operations

Lowering rise.codegen to loop: Matrix Multiplication

```
1 func @mm(%out: memref<1024x1024xf32>, %A: memref<1024x1024xf32>, %B: memref<1024x1024xf32>) {
2   %cst0 = constant 0.000000e+00 : f32
3   %c0 = constant 0 : index
4   %c1024 = constant 1024 : index
5   %c1 = constant 1 : index
6   loop.for %i = %c0 to %c1024 step %c1 {
7
8
9     loop.for %j = %c0 to %c1024 step %c1 {
10
11
12
13
14       store %cst0, %out[%j, %i] : memref<1024x1024xf32>
15       loop.for %k = %c0 to %c1024 step %c1 {
16         %0 = load %A[%k, %i] : memref<1024x1024xf32>
17         %1 = load %B[%k, %j] : memref<1024x1024xf32>
18         %2 = load %out[%j, %i] : memref<1024x1024xf32>
19
20         %3 = mulf %0, %1 : f32
21         %4 = addf %3, %2 : f32
22         store %4, %out[%j, %i] : memref<1024x1024xf32>
23       }
24     }
25   }
26   return
27 }
```

Loop

Lowering RISE to affine

```
1 func @mm(%out:memref<1024x1024xf32>,%A:memref<1024x1024xf32>,  
2 %B:memref<1024x1024xf32>) {  
3   %cst0 = constant 0.000000e+00 : f32  
4   %c0 = constant 0 : index  
5  
6  
7   affine.for %i = 0 to 1024 {  
8     affine.for %j = 0 to 1024 {  
9       affine.store %cst0, %out[%j, %i] : memref<1024x1024xf32>  
10      affine.for %k = 0 to 1024 {  
11        %0 = affine.load %arg1[%k, %i] : memref<1024x1024xf32>  
12        %1 = affine.load %arg2[%k, %j] : memref<1024x1024xf32>  
13        %2 = affine.load %arg0[%j, %i] : memref<1024x1024xf32>  
14        %3 = mulf %0, %1 : f32  
15        %4 = addf %3, %2 : f32  
16        affine.store %4, %arg0[%j, %i] : memref<1024x1024xf32>  
17      }  
18    }  
19  }  
20 }  
21 return
```

Affine

- Similar to the loop lowering presented before
- Lowering target specified using Attributes
 - `rise.reduceSeq` {to = "affine"}
- We can always use affine loops with our patterns
- Enables usage of polyhedral optimizations

Lowering RISE to affine vs. loop

```
1 func @mm(%out:memref<1024×1024xf32>,%A:memref<1024×1024xf32>,  
2 %B:memref<1024×1024xf32>) {  
3   %cst0 = constant 0.000000e+00 : f32  
4   %c0 = constant 0 : index  
5  
6  
7   affine.for %i = 0 to 1024 {  
8     affine.for %j = 0 to 1024 {  
9       affine.store %cst0, %out[%j, %i] : memref<1024×1024xf32>  
10      affine.for %k = 0 to 1024 {  
11        %0 = affine.load %arg1[%k, %i] : memref<1024×1024xf32>  
12        %1 = affine.load %arg2[%k, %j] : memref<1024×1024xf32>  
13        %2 = affine.load %arg0[%j, %i] : memref<1024×1024xf32>  
14        %3 = mulf %0, %1 : f32  
15        %4 = addf %3, %2 : f32  
16        affine.store %4, %arg0[%j, %i] : memref<1024×1024xf32>  
17      }  
18    }  
19  }  
20  return  
21 }
```

Affine

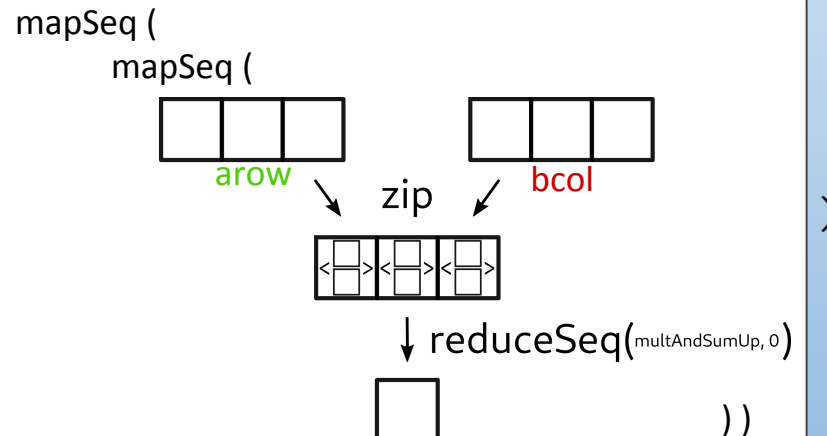
```
1 func @mm(%out: memref<1024×1024xf32>, %A: memref<1024×1024xf32>,  
2 %B: memref<1024×1024xf32>) {  
3   %cst0 = constant 0.000000e+00 : f32  
4   %c0 = constant 0 : index  
5   %c1024 = constant 1024 : index  
6   %c1 = constant 1 : index  
7   loop.for %i = %c0 to %c1024 step %c1 {  
8     loop.for %j = %c0 to %c1024 step %c1 {  
9       store %cst0, %out[%j, %i] : memref<1024×1024xf32>  
10      loop.for %k = %c0 to %c1024 step %c1 {  
11        %0 = load %A[%k, %i] : memref<1024×1024xf32>  
12        %1 = load %B[%k, %j] : memref<1024×1024xf32>  
13        %2 = load %out[%j, %i] : memref<1024×1024xf32>  
14        %3 = mulf %0, %1 : f32  
15        %4 = addf %3, %2 : f32  
16        store %4, %out[%j, %i] : memref<1024×1024xf32>  
17      }  
18    }  
19  }  
20  return  
21 }
```

Loop

Lowering RISE to library calls



match_for(



Lowering RISE to library calls



match_for(


matchSuccess

mapSeq (

mapSeq (



zip



reduceSeq(multAndSumUp, 0)




))

```
1 func @mm(%out:memref<1024x1024xf32>,%A:memref<1024x1024xf32>,%B:memref<1024x1024xf32>) {
2   %CblasRowMajor = constant 101 : i32
3   %CblasNoTrans = constant 111 : i32
4   %M = constant 1024 : i32
5   %N = constant 1024 : i32
6   %K = constant 1024 : i32
7   %LDA = constant 1024 : i32
8   %LDB = constant 1024 : i32
9   %LDC = constant 1024 : i32
10  %alpha = constant 1.0 : f32
11  %beta = constant 1.0 : f32
12  call @cblas_sgemm_wrapper(%CblasRowMajor, %CblasNoTrans, %CblasNoTrans, %M, %N, %K, %alpha, %A, %LDA, %B, %LDB, %beta, %out, %LDC) :
13    (i32, i32, i32, i32, i32, i32, f32, memref<1024x1024xf32>, i32, memref<1024x1024xf32>, i32, f32, memref<1024x1024xf32>, i32) → ( )
14  return
15 }
```

State of our implementation

- Lambda calculus foundations implemented
- Basic set of patterns implemented
- Basic lowering system implemented
- Lowering to Library Calls quite ad-hoc
- Preliminary experiments shows performance as expected



RISE in MLIR: Summary

Today

- Lambda-Calculus + Composable Patterns in MLIR
- Ability to embed arbitrary code from other dialects
- Different lowering approaches

Future directions:

- Explore other lowering approaches (e.g. GPU, OpenMP dialects)
- Introduce symbolic sizes via dependent types
- Rewriting enables composable and reusable high-level transformations
New language for controlling rewriting (ELEVATE - <https://elevate-lang.org>)

RISE in MLIR

A functional Pattern-based Dialect

We are Open Source!

<https://rise-lang.org/mlir>

<https://github.com/rise-lang/mlir>

Martin Lücke | **Michel Steuwer** | **Aaron Smith**



THE UNIVERSITY
of EDINBURGH



University
of Glasgow